



# EECS4030: Computer Architecture

## Instructions: Language of the Computer (I)

Prof. Chung-Ta King

Department of Computer Science

National Tsing Hua University, Taiwan

(Adapted from textbook slides <https://www.elsevier.com/books-and-journals/book-companion/9780128122754/lecture-slides>)



國立清華大學

National Tsing Hua University

# Introduction

- To command a computer, you must speak its language
- Most computers understand high-level languages (HLL), e.g., C, C++, Java, with suitable compilers

```
foo(int v[], int k) {  
    int temp;  
    temp = v[k];  
    v[k] = v[k] + v[k+1];  
    v[k+1] = temp;  
}
```

**Variables** (memory locations storing data)

**operands** of operators

**Operators**

**operations** on operands

- How about language to a computer's hardware?
  - *Assembly language* → *machine instructions*
  - Very close to control signals to control hardware operations





# Outline

- Languages of computers (Sec. 2.1)
- Operations of computer hardware (Sec. 2.2)
- Operands of computer hardware (Sec. 2.3, 2.4)
  - Registers, memory, signed and unsigned numbers
- Representing instructions (Sec. 2.5)
- More operations: logic, decision making (Sec. 2.6, 2.7)
- Supporting procedures in hardware (Sec. 2.8)
- More operand representations (Sec. 2.9, 2.10)
  - Characters, wide immediates and addresses
- Parallelism and instructions (Sec. 2.11)
- Translating and starting a program (Sec. 2.12)





# Outline

- A C sort example to put it all together (Sec. 2.13)
- Arrays versus pointers (Sec. 2.14)
- Compiling C and interpreting Java (Sec. 2.15)
- Other ISAs: MIPS, x86 (Sec. 2.16, 2.17)
- The rest of RISC-V (Sec. 2.18)
- Fallacies and pitfalls (Sec. 2.19)



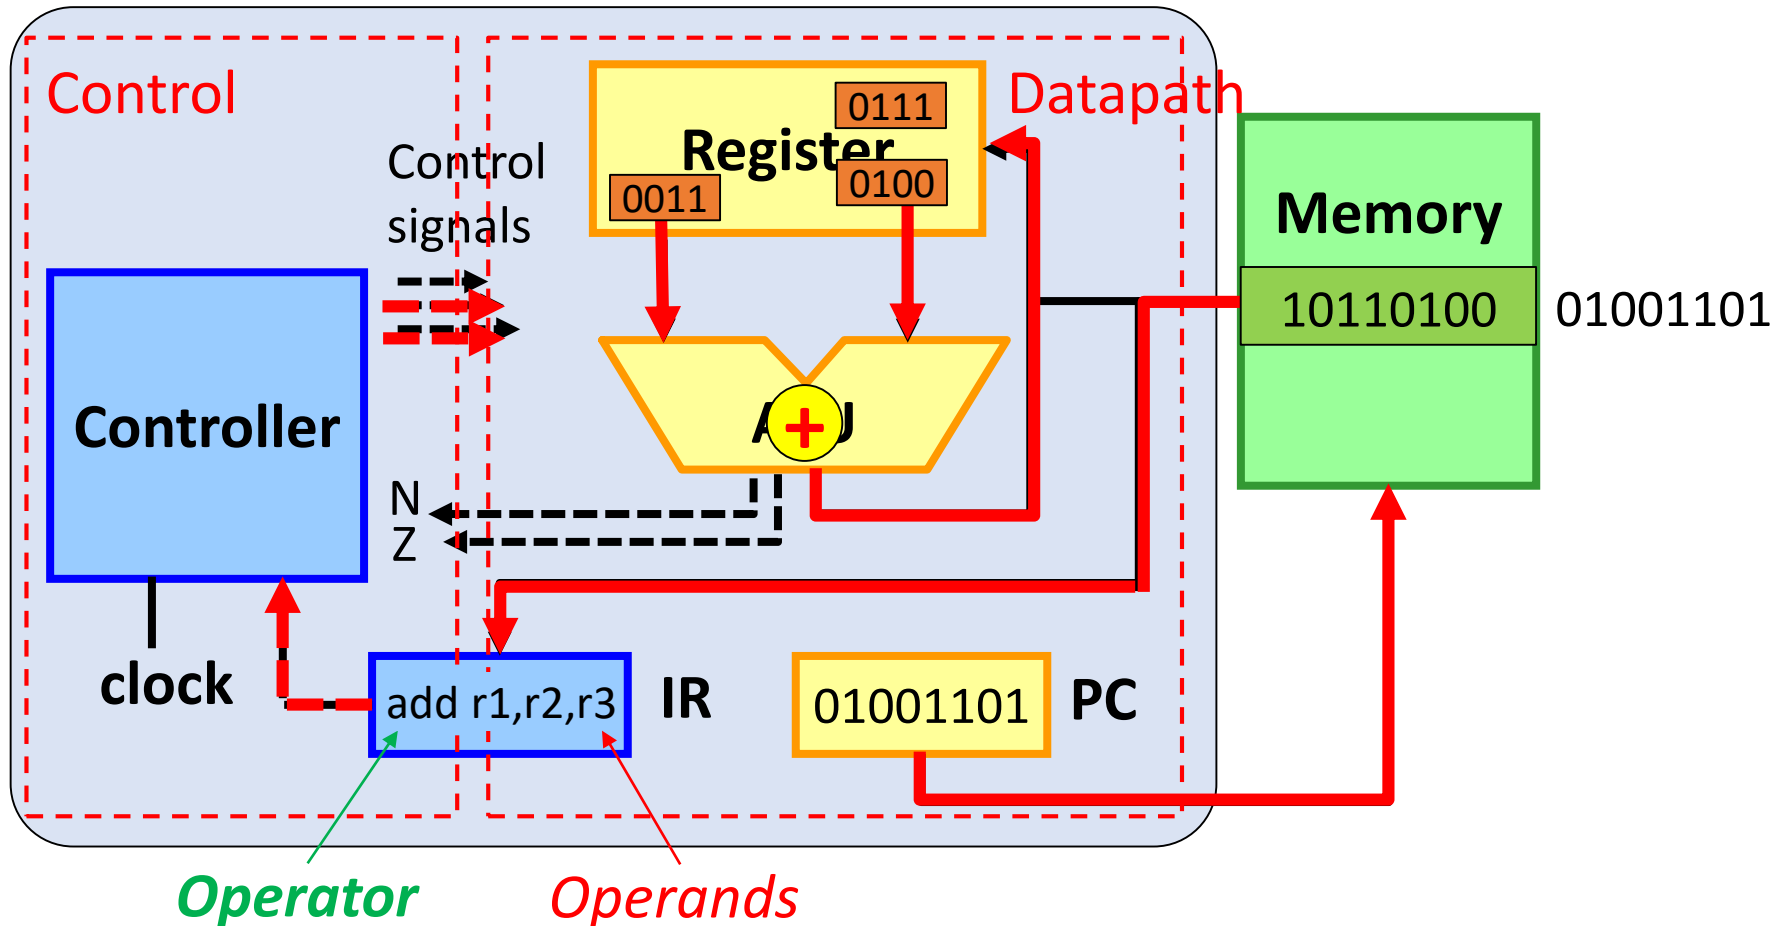


# Outline

- Languages of computers (Sec. 2.1)
- Operations of computer hardware (Sec. 2.2)
- Operands of computer hardware (Sec. 2.3, 2.4)
  - Registers, memory, signed and unsigned numbers
- Representing instructions (Sec. 2.5)
- More operations: logic, decision making (Sec. 2.6, 2.7)
- Supporting procedures in hardware (Sec. 2.8)
- More operand representations (Sec. 2.9, 2.10)
  - Characters, wide immediates and addresses
- Parallelism and instructions (Sec. 2.11)
- Translating and starting a program (Sec. 2.12)



# (Command) Language to Hardware



**Instruction:** a sentence of the language, do something (on some data)





# Language to Computer Hardware

- Different computers have different languages (instruction sets)
  - But with many aspects in common
  - Because all computers are constructed from hardware technologies based on similar underlying principles
  - Because there are a few basic operations that all computers must provide
    - Makes it easy to build the hardware and the compiler while maximizing performance and minimizing cost and energy
    - Combinations of these basic operations allow a computer to run all sorts of programs written in various HLLs

**What are the most basic operations and how to specify them?**



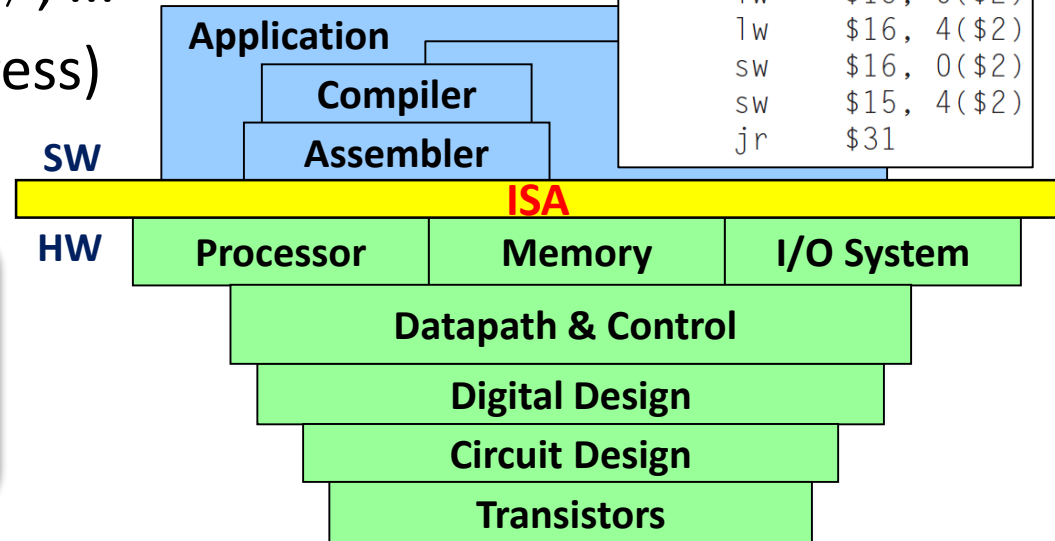
# Essentially, We Are Asking ...



- What must be specified in an instruction set architecture (ISA)?
  - “ISA encompasses all the information necessary to write a machine language program that will run correctly”
  - Including:
    - operations**: +, −, \*, /, ...
    - operands** (data, address)
    - flow control**

```
swap:
    multi $2, $5, 4
    add   $2, $4, $2
    lw    $15, 0($2)
    lw    $16, 4($2)
    sw    $16, 0($2)
    sw    $15, 4($2)
    jr    $31
```

What are included in ISA and how are they specified have profound effects on computer design!







# Components of an ISA

## Operands: storage and data

- Organization of programmable storage:
  - **Registers**: # of bits, # of registers, their roles, ...
  - **Memory**: **addressing**, access methods, ...
- Data types and data structures
  - Integers, floating-point, text, signed number, ...

## Operations: on data and on execution flow

- Instruction set (or operation code)
  - ALU, control transfer, exceptional handling

## Instruction specification:

- Instruction format and encoding





# Use 64-bit RISC-V (RV64) as an Example

- Developed at UC Berkeley as open ISA
  - Now managed by the RISC-V Foundation ([riscv.org](https://riscv.org))
- Typical of many modern ISAs
- Similar ISAs have a large share of embedded core market
  - Applications in consumer electronics, network/storage equipment, cameras, printers, ...
- We will not only describe what the RISC-V ISA is
- But also study the **principles** behind the design
- *And examine some different design **options** to appreciate why RISC-V chooses the specific option*





# Outline

- Languages of computers (Sec. 2.1)
- **Operations of computer hardware (Sec. 2.2)**
- Operands of computer hardware (Sec. 2.3, 2.4)
  - Registers, memory, signed and unsigned numbers
- Representing instructions (Sec. 2.5)
- More operations: logic, decision making (Sec. 2.6, 2.7)
- Supporting procedures in hardware (Sec. 2.8)
- More operand representations (Sec. 2.9, 2.10)
  - Characters, wide immediates and addresses
- Parallelism and instructions (Sec. 2.11)
- Translating and starting a program (Sec. 2.12)



# What **Basic** Operations for Hardware?

1. Operations to work on data
  - Arithmetic (+, -, \*, /), logic (&, |, ~)  
→ similar to a function, e.g.  $z=f(x,y)$ , operating on values
  - 2-operand ( $a + b$ ), 1-operand ( $++a$ )
2. Operations to move and copy data between storage locations
  - As storage, original location still retains the data, until explicitly overwritten
3. Operations to change execution flow
  - *Unconditional*: goto, jump to new location to execute
  - *Conditional*: if-then-else, jump if condition is true
  - *Procedural*:  $foo()$ , similar to goto but need to come back



# Let's Start with Arithmetic Operations

- The most basic operation for any computer hardware is arithmetic operation on numbers, e.g.,

$$a = b + c$$

- Basic RISC-V arithmetic instructions are specified as:

**add a, b, c                      # a ← b + c**

- All RISC-V arithmetic instructions have this form

- **3-address**: **a**, **b**, **c** are storage locations that hold data, but **c** may also be a value

*Why storage concept indispensable?*

## **Design Principle 1: *Simplicity favors regularity***

- Regularity makes implementation simpler
- Simplicity enables higher performance at lower cost

**vs. Flexibility**



# Arithmetic Instruction Example

- C code:

```
f = (g + h) - (i + j);
```

- Compiled RISC-V (pseudo) code:

```
add t0, g, h    # temp t0 ← g + h  
add t1, i, j    # temp t1 ← i + j  
sub f, t0, t1   # f ← t0 - t1
```

- **t0** and **t1** are temporary variables, i.e., storage locations, normally created and allocated by the compiler



# Why 3-Address Format: add a,b,c?

- Why not 2-address format?

**add b, c                      # b  $\leftarrow$  b + c**

- Still regular, and one address less to specify and encode!  
→ Many 16-bit processors use this format

- How about 0-address format?

**add**

- Where to get the operands?  
→ must be specified implicitly so that the computer hardware knows where to get, e.g. in a stack

**add    #stack[top]  $\leftarrow$  stack[top] + stack[next]**

Tradeoff between flexibility and complexity





# Outline

- Languages of computers (Sec. 2.1)
- Operations of computer hardware (Sec. 2.2)
- **Operands of computer hardware (Sec. 2.3, 2.4)**
  - Registers, memory, signed and unsigned numbers
- Representing instructions (Sec. 2.5)
- More operations: logic, decision making (Sec. 2.6, 2.7)
- Supporting procedures in hardware (Sec. 2.8)
- More operand representations (Sec. 2.9, 2.10)
  - Characters, wide immediates and addresses
- Parallelism and instructions (Sec. 2.11)
- Translating and starting a program (Sec. 2.12)





# Operands to an Operation

- Operands may be the value itself (*immediates*) 3  
    `add a, b, 3`  
or the storage location holding the value  
    `add a, b, c`2
- The storages that can hold values may be in *memory*  
or in *registers* 1
  - Registers are a small number of storage locations built directly into CPU hardware; memory refers to storage locations in the main memory → need to specify *addresses*
  - How about *disk*? *Why not disk locations as operands?*
- RISC-V arithmetic instructions must use **register operands!**



# Why Restricted to Register Operands?

- Why not allow memory operands, too?

**add M[1000], M[2019], M[2020]**

## **Design Principle 2: *Smaller is faster***

- Register vs. main memory (millions of locations)
  - Registers are inside CPU → same clock as CPU, connect directly to arithmetic units, short and fast signals
  - Registers are small → easy to activate to access data, short and fast signals
  - Registers are small → fewer number of bits to specify and encode → smaller instructions
  - Suitable for holding frequently accessed data



# How to Specify (1) Register Operands?

- Must know RISC-V register organization!
- RISC-V has a  $32 \times 64$ -bit *register file*
  - 32 registers specified as **x0** to **x31**, each with a **5-bit** addr

**doubleword**

x0 constant 0	x10 function arguments/results x11
x1 return address	
x2 stack pointer	x12 ⋮ function arguments x17
x3 global pointer	
x4 thread pointer	x18 ⋮ saved registers x27
x5 ⋮ temporaries x7	
x8 frame pointer	
x9 saved registers	x28 ⋮ temporaries x31

**Why not 64,  
or 16?**



# Register Operand Example

- How to do the following C statement?

$$f = (g + h) - (i + j);$$

- Assume that  $f, g, h, i, j$  are in  $x19, x20, x21, x22, x23$
- Use intermediate temporary registers  $x5, x6$

<code>add x5, x20, x21</code>	<code># x5 ← g + h</code>
<code>add x6, x22, x23</code>	<code># x6 ← i + j</code>
<code>sub x19, x5, x6</code>	<code># f ← x5 - x6</code>

IC↑  
IPC↓

- If data are in the memory, they must be loaded into the registers first before arithmetic operations can operate on them; after operations, they need to be stored back



Use *data transfer instructions* (`ld`, `sd`, ...) discussed later



# How to Specify (2) Memory Operands?

- Must know RISC-V memory organization!
  - Main memory is organized as a 1D array of locations
  - Each location can hold one byte of data and is identified by an *address* → memory is *byte addressed*, i.e. each address identifies an 8-bit byte of data, ranging from 0 to  $2^{61}$
- A nature way of specifying memory operands in RISC-V is like addressing an array, e.g., **A[100]**
  - **A** is the **base address** and 100 is the **offset** and memory address of **A[100]** is (**A** + 100)
  - In RISC-V, memory operand is specified as: **100(x22)**, where **x22** contains base addr. and **100** is offset (in bytes)
  - A memory operand normally gets a doubleword (64 bits), and thus word addresses differ by 8



# Memory Operand Example

- C code:

**A[9] = h + A[8];**

- **h** in **x21**, *base address* of **A** in **x22**
- **A[8]**: the 8<sup>th</sup> doubleword in array **A**[], which is at the 64<sup>th</sup> byte from start of **A**[]

- Compiled RISC-V code:

**ld x9, 64(x22)**

*load doubleword*

**add x9, x21, x9**

**sd x9, 72(x22)**

*store doubleword*

offset

base register


$2^{64} - 1$

A[9] = 1096

A[8] = 1088


A = 1024

0



# Registers and Memory Size

## RISC-V CPU



5 bits for register ID

64-bit address

0x FFFF FFFF FFFF FFFF

...

0x 0000 0000 0000 0008

0x 0000 0000 0000 0007

0x 0000 0000 0000 0006

0x 0000 0000 0000 0005

0x 0000 0000 0000 0004

0x 0000 0000 0000 0003

0x 0000 0000 0000 0002

0x 0000 0000 0000 0001

0x 0000 0000 0000 0000

data

Memory



double  
word  
=64 bits

1 byte = 8 bits

**Note:** actual physical address bits may be fewer than 64, e.g., 42 for Intel x86\_64

(Prof. Jing-Jia Liou)

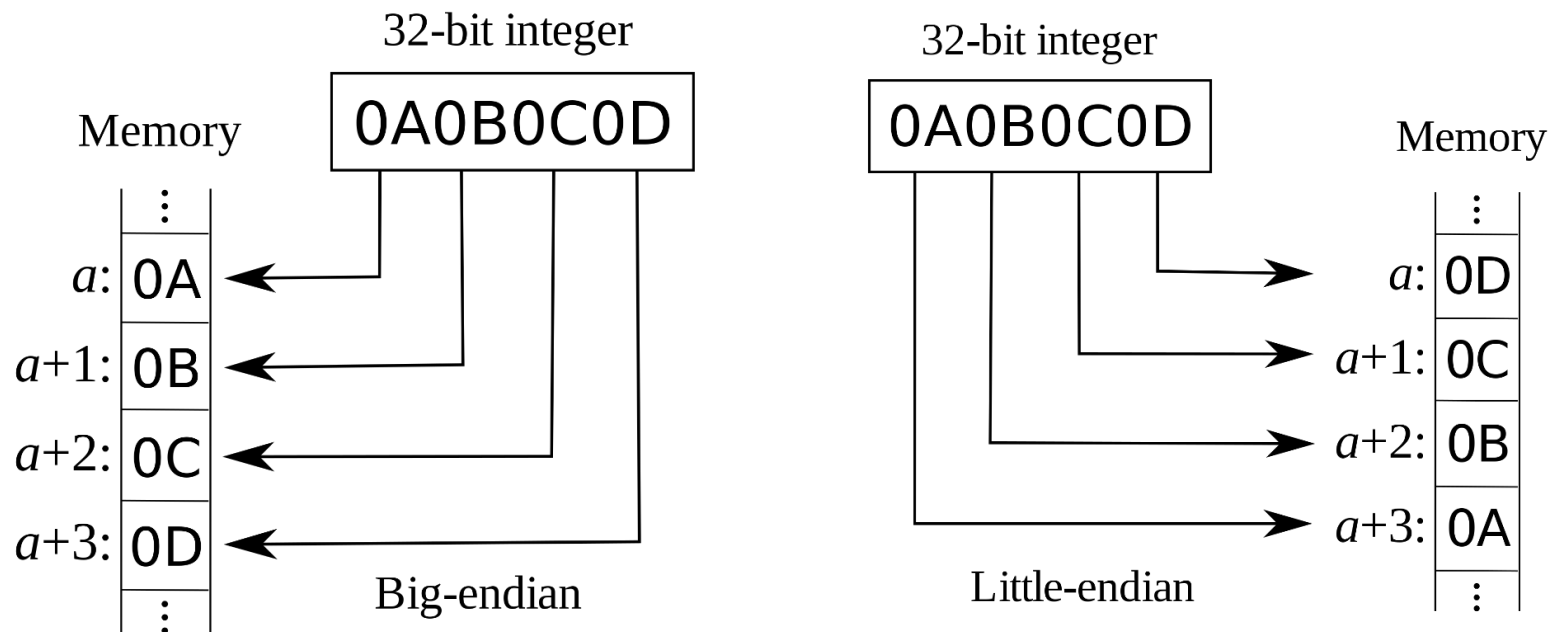


國立清華大學

National Tsing Hua University

# More about RISC-V Memory Organization

- RISC-V is *Little Endian*
  - Least-significant byte at least address of a word
  - c.f. Big Endian: most-significant byte at least address



(By R. S. Shaw - Own work, Public Domain, <https://commons.wikimedia.org/w/index.php?curid=2974661>)

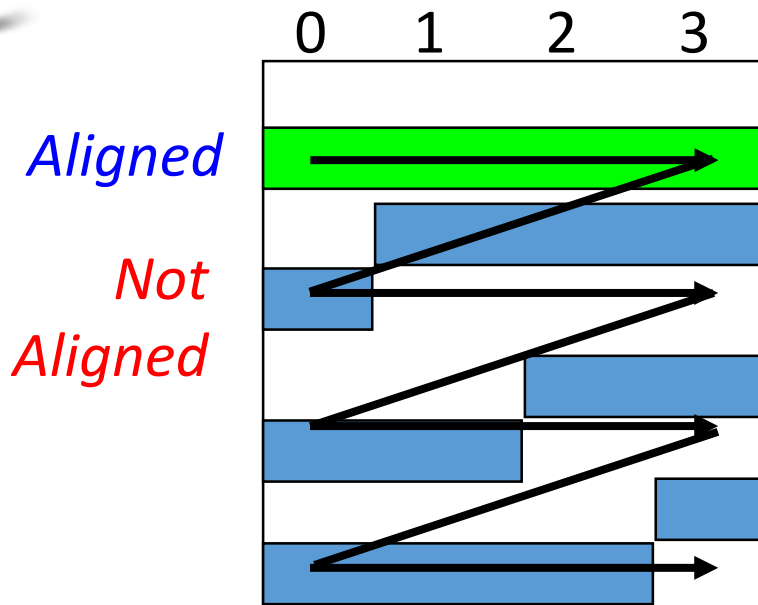




# More about RISC-V Memory Organization

- RISC-V does not require words to be **aligned** in memory, i.e., at an address that is a multiple of 4 (for a word) or 8 (for a doubleword)

- Unlike some other ISAs; but misaligned memory accesses may be very very slow



Not  
Aligned

Aligned


$A[9] = 1096$

$A[8] = 1088$

$A = 1024$

0

$2^{64} - 1$





# Registers vs. Memory

- Registers are faster to access than memory
  - Also, operating on memory data requires loads and stores
    - More instructions to be executed, larger code size →  $IC \uparrow$
- Compilers should use registers for variables as much as possible →  $IPC \downarrow$ 
  - Only *spill* to memory for less frequently used variables, if more variables than available registers are used
  - Register optimization is important!
- Why not keep all variables in memory?
  - Smaller is faster: registers are faster than memory
  - Registers are more versatile:
    - Register file can read 2 and write 1 registers per instruction
    - Memory read/write 1 operand per instruction, no operation





# How about (3) Immediate Operands?

- Small constants used frequently (50% of operands)  
e.g.,  $A = B + 4$ ; if ( $C > 1$ ) ... ;

- Put 'typical constants' in memory and load them

```
ld    x9, AddrConstant4(x3)
```

```
add   x22, x22, x9
```

Too slow

- Alternative: constant data specified in the instruction

```
addi   x22, x22, 4
```

Immediate operand

## Design Principle 3: *Make the common case fast*

- Small constants are common
- Immediate operand avoids a load instruction

Why define a new instruction, **addi**? Why not reuse **add**?



# Handling Constant Zero and Negatives

- The number zero (0) is so common in code that RISC-V hardwired register 0 (**x0**) to 0
  - Cannot be overwritten, e.g., **addi x0, x0, 5** (no effect)
  - Useful for other common operations, too, e.g., move between registers: **add x22, x20, x0**
    - Thus, there is no need to implement a move instruction in hardware → simpler CPU design and implementation
- No need for subtract immediate instruction either
  - Just use a negative constant: **addi x22, x22, -8**

*Time to review signed and unsigned numbers*



# Represent **Unsigned** Binary Integers

- Given an  $n$ -bit number

$$x = x_{n-1}2^{n-1} + x_{n-2}2^{n-2} + \dots + x_12^1 + x_02^0$$

- Use all the bits to represent **values**
- Example: (32 bits)

*Can get the max value*

0000 0000 0000 0000 0000 0000 0000 1011<sub>2</sub>

$$= 0 \times 2^{31} + \dots + 1 \times 2^3 + 0 \times 2^2 + 1 \times 2^1 + 1 \times 2^0$$

$$= 0 + \dots + 8 + 0 + 2 + 1 = 11_{10}$$

- Range: 0 to  $+2^n - 1$ 
  - Using 32 bits: 0 to +4,294,967,295
  - Using 64 bits: 0 to +18,446,774,073,709,551,615



# Signed Binary Integers (2's-Complement)

- Given an  $n$ -bit number

$$X = -x_{n-1}2^{n-1} + x_{n-2}2^{n-2} + \dots + x_12^1 + x_02^0$$

- Use bit 31/63 to represent **sign** (*sign bit*)
  - 1 for negative and 0 for non-negative numbers
- Signed negation: complement ( $1 \rightarrow 0, 0 \rightarrow 1$ ) & add 1
  - Non-negative numbers have the same representation
  - $-(-2^{n-1})$  can't be represented

- Example: (32 bits)

$$1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1100_2$$

$$= -1 \times 2^{31} + 1 \times 2^{30} + \dots + 1 \times 2^2 + 0 \times 2^1 + 0 \times 2^0$$

$$= -2,147,483,648 + 2,147,483,644 = -4_{10}$$




# 2's-Complement Signed Integers

- Range:  $-2^{n-1}$  to  $+2^{n-1} - 1$ 
  - Using 32 bits:  $-2,147,483,648$  to  $+2,147,483,647$
  - Using 64 bits:  $-9,223,372,036,854,775,808$  ( $1000 \dots 0000_2$ ) to  $9,223,372,036,854,775,807$  ( $0111 \ 1111 \dots 1111_2$ )
  - 0:  $0000 \ 0000 \dots 0000_2$
  - -1:  $1111 \ 1111 \dots 1111_2$  (least negative number)
  - Most-negative number ( $1000 \dots 0000_2$ ) does not have a corresponding positive number

**Why 2's  
complement?**

*One adder, no distinction*

*SW interprets meaning*



*Signed/unsigned +/-  
treated the same!*



# Sign Extension

- Representing a number using more bits while preserving the numeric value
- How to do? replicate the sign bit to the left
  - c.f. unsigned values: extend with 0s
- Examples: 8-bit to 16-bit

Sign?

– +2: 0000 0010 → 0000 0000 0000 0010

– -2: 1111 1110 → 1111 1111 1111 1110

1b x22, 0 (x21)

FFFFFF F7

- In RISC-V instruction set 1bu x22, 0 (x21)

000000 F7

- 1b: load byte (sign-extend loaded byte)
- 1bu: load byte unsigned (zero-extend loaded byte)







# Outline

- Languages of computers (Sec. 2.1)
- Operations of computer hardware (Sec. 2.2)
- Operands of computer hardware (Sec. 2.3, 2.4)
  - Registers, memory, signed and unsigned numbers
- **Representing instructions (Sec. 2.5)**
- More operations: logic, decision making (Sec. 2.6, 2.7)
- Supporting procedures in hardware (Sec. 2.8)
- More operand representations (Sec. 2.9, 2.10)
  - Characters, wide immediates and addresses
- Parallelism and instructions (Sec. 2.11)
- Translating and starting a program (Sec. 2.12)



# How to Specify RISC-V Instructions?

- Now that we have learned *operations* and *operands* of computer hardware, let us study how to put them together to form computer instructions
  - Remember: Computer only understands 1s and 0s, so “add x21, x22, x23” is meaningless to hardware
  - Instructions to instruct computer hardware how to operate should also be encoded in binary, called *machine code*
- RISC-V’s strategies for instruction encoding
  - Encode every instruction in **32-bit** words, divided into *fields*
  - Small number of formats encoding *operation code* (*opcode*), register numbers, ...
- Let’s look at three of them first

**Regularity!**

*Why not just 1 format?*



# RISC-V R-format Instructions

funct7	rs2	rs1	funct3	rd	opcode
7 bits	5 bits	5 bits	3 bits	5 bits	7 bits

- For register to register operations
- Instruction fields
  - opcode: operation code
  - rd: destination register number
  - funct3: 3-bit function code (additional opcode)
  - rs1: the first source register number
  - rs2: the second source register number
  - funct7: 7-bit function code (additional opcode)

*Why not combine  
op and funct?*



# R-format Example

funct7	rs2	rs1	funct3	rd	opcode
7 bits	5 bits	5 bits	3 bits	5 bits	7 bits

**add x9, x20, x21**

0	x21	x20	0	x9	add
---	-----	-----	---	----	-----

0	21	20	0	9	51
---	----	----	---	---	----

0000000	10101	10100	000	01001	0110011
---------	-------	-------	-----	-------	---------

0000 0001 0101 1010 0000 0100 1011 0011<sub>two</sub> = 015A04B3<sub>16</sub>



# RISC-V I-format Instructions

immediate	rs1	funct3	rd	opcode
12 bits	5 bits	3 bits	5 bits	7 bits

- Immediate arithmetic and load instructions
  - rs1: source or base address register number
  - immediate: constant operand, or byte offset to base addr
    - 2's-complement, sign extended,  $-2^{11}$  to  $2^{11}-1$

**Design Principle 4:** *Good design demands good compromises*

- Different formats complicate decoding, but allow 32-bit instructions uniformly
- Keep formats as similar as possible (only 1 field differs)



# I-Format Example 1

- RISC-V instruction:

**addi      x9 , x22 , -51**

- opcode = 19 (look up in RISC-V Manual)
- rs1 = 22 (source operand register number)
- rd = 9 (destination register number)
- immediate = -51 (by default, this is decimal)

- Decimal representation:

-51	22	0	9	19
-----	----	---	---	----

- Binary representation:

1111 1100 1101	10110	000	01001	0010011
----------------	-------	-----	-------	---------



# I-Format Example 2

- RISC-V instruction:

**ld x9, 120(x22)**

- opcode = 3 (look up in RISC-V Manual)
- rs1 = 22 (base address register number)
- rd = 9 (destination register number)
- immediate = 120 (offset)

- Decimal representation:

120	22	3	9	3
-----	----	---	---	---

- Binary representation:

0000 0111 1000	10110	011	01001	0000011
----------------	-------	-----	-------	---------



# RISC-V S-format Instructions

imm[11:5]	rs2	rs1	funct3	imm[4:0]	opcode
-----------	-----	-----	--------	----------	--------

7 bits

5 bits

5 bits

3 bits

5 bits

7 bits

- Different immediate format for store instructions
  - rs1: base address register number
  - rs2: source operand register number
  - immediate: offset added to base address
    - Split so that rs1 and rs2 fields always in the same place

- RISC-V instruction:

**sd      x9, 120 (x22)**

120(1)	9	22	3	120(2)	35
0000011	01001	10110	011	11000	0100011







# Key Computer Design Principles

- Computer instructions represented as numbers
- Program instructions can be stored at the same place as data (in memory) and be manipulated in the same way as data → **stored-program computer**
  - Instructions and data look the same, just 0s and 1s
  - Differ only in how you interpret the bits
- One consequence: everything addressed
  - Everything has a memory address: instructions, data
    - Both branches and jumps use these
  - One register keeps address of the instruction being executed: “**Program Counter**” (PC)
  - Register can hold any 64-bit value, (signed) int, address, etc.

**Key to  
programmable  
general-purpose**



# Stored-Program Computers

- Consequence #2:  
programs can operate on programs
  - e.g., compilers, linkers, ...
  - The latter is data to the former
- Consequence #3:  
*binary compatibility*
  - Programs can be shipped as files of binary numbers and executed on any computers with a compatible ISA
- Consequence #4:  
fast context switches of processes

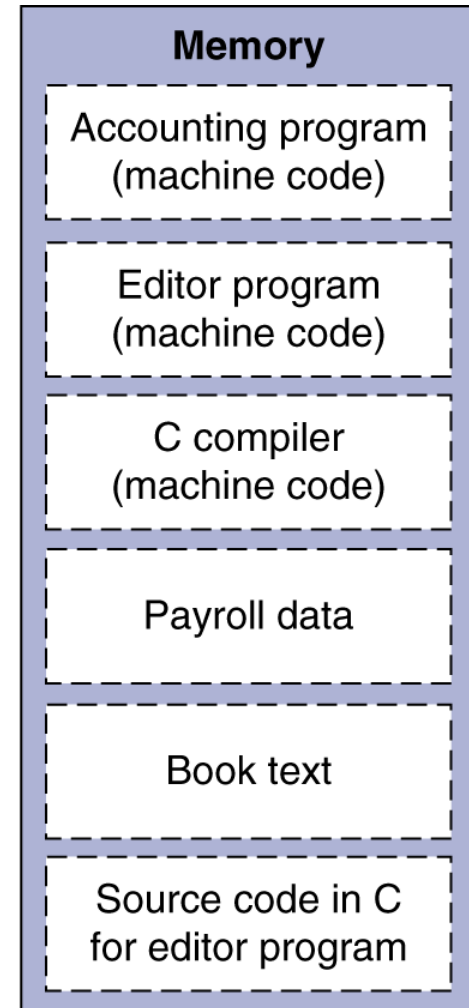

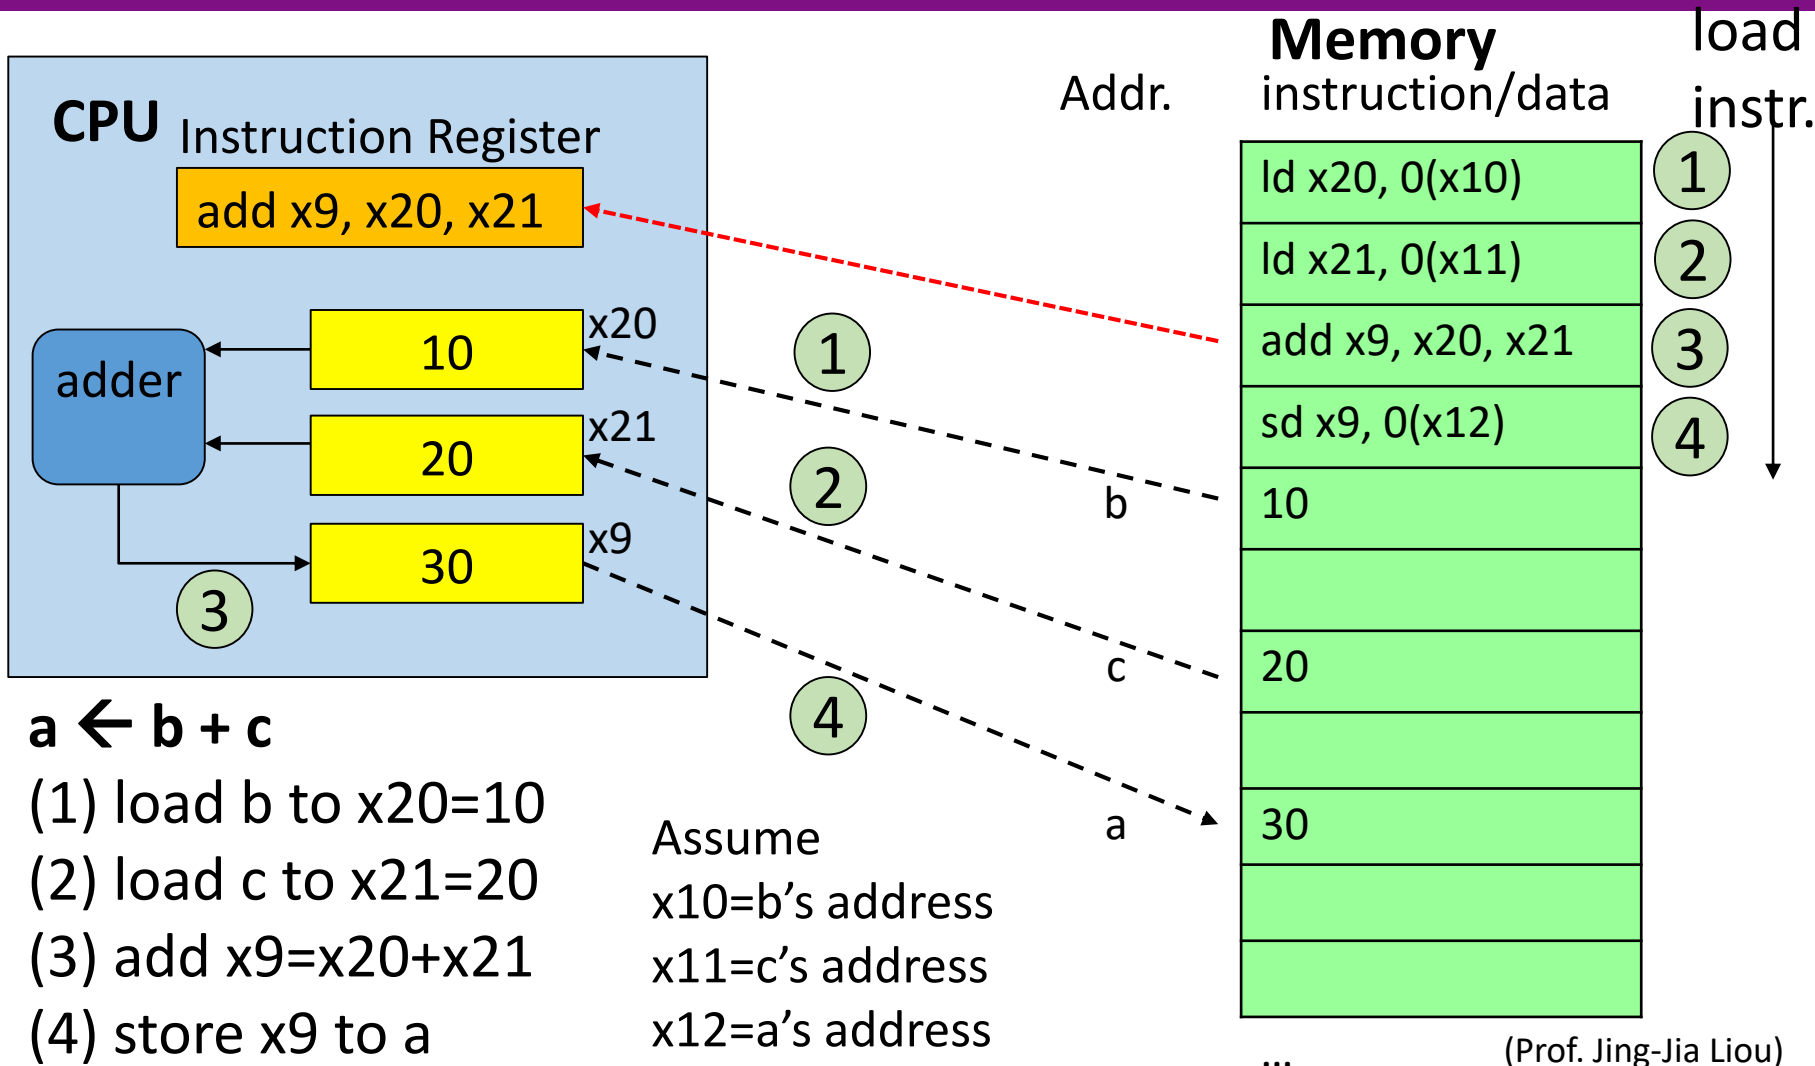


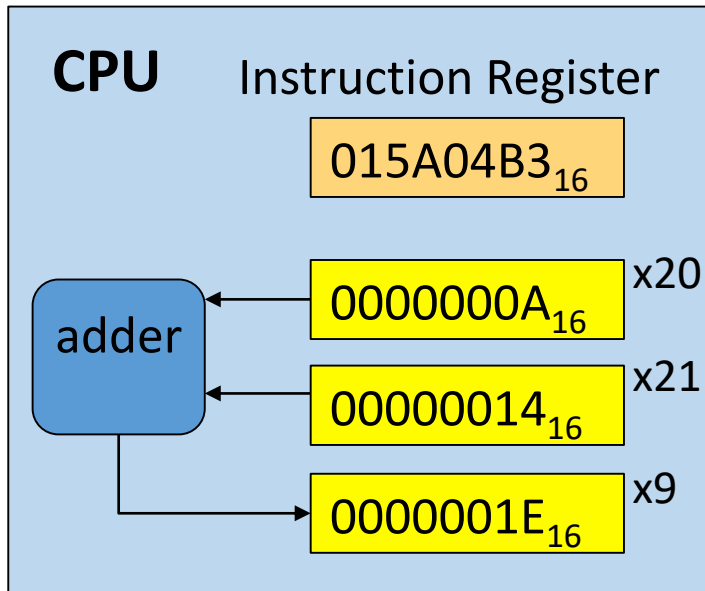
Fig. 2.7



# Stored Program with Memory Access



# Instruction and Data Encoding (in Hex)



$$a \leftarrow b + c$$

- (1) load b to x20=10
- (2) load c to x21=20
- (3) add x9=x20+x21
- (4) store x9 to a

Assume  
 x10=b's address  
 x11=c's address  
 x12=a's address

Memory	
address	instruction/data
	00053A03 <sub>16</sub>
	0005BA83 <sub>16</sub>
	015A04B3 <sub>16</sub>
	00963023 <sub>16</sub>
b	0000000A <sub>16</sub>
c	00000014 <sub>16</sub>
a	0000001E <sub>16</sub>
...	

(Prof. Jing-Jia Liou)





# Outline

- Languages of computers (Sec. 2.1)
- Operations of computer hardware (Sec. 2.2)
- Operands of computer hardware (Sec. 2.3, 2.4)
  - Registers, memory, signed and unsigned numbers
- Representing instructions (Sec. 2.5)
- **More operations: logic, decision making (Sec. 2.6, 2.7)**
- Supporting procedures in hardware (Sec. 2.8)
- More operand representations (Sec. 2.9, 2.10)
  - Characters, wide immediates and addresses
- Parallelism and instructions (Sec. 2.11)
- Translating and starting a program (Sec. 2.12)





# Bitwise Operations

- Up until now, we have seen arithmetic (**add**, **addi**) and memory access (**ld**, **sd**)
  - All of these instructions view contents of register as a single quantity (such as a signed or unsigned integer)
- **New perspective:** View contents of register as 64 individual bits rather than as a single 64-bit number
  - Since registers are composed of 64 bits, we may want to access individual bits rather than the whole
- Introduce two new classes of instructions:
  - Logical operators
  - Shift instructions



# Logical and Shift Operations

- Instructions for **bitwise** manipulation
  - Useful for extracting and inserting groups of bits in a word

Operation	C	Java	RISC-V
Shift left	<<	<<	sll, slli
Shift right	>>	>>>	srl, srli
Shift right arith.	>>	>>	sra, srai
Bitwise AND	&	&	and, andi
Bitwise OR			or, ori
Bitwise XOR	^	^	xor, xori
Bitwise NOT	~	~	

Fig. 2.8



# Shift Operations

- sll: shift left logical (R-type)
  - Shift rs1 left for rs2 bits and fill rs1 with 0 bits
  - Shift left by  $i$  bits = multiplies by  $2^i$  (faster than multiply)

funct7	rs2	rs1	funct3	rd	opcode
7 bits	5 bits	5 bits	3 bits	5 bits	7 bits

- Example: **sll x9, x20, x21**

0	x21	x20	1	x9	sll
0000000	10101	10100	001	01001	0110011

- **x20** = 0000 ... 0001; **x21** = 0000 ... 0011 = 3
- After instruction executed, **x9** will contain 0000 ... 1000





# Shift Operations

- srl: shift right logical (R-type)
  - Shift rs1 right for rs2 bits and fill rs1 with 0 bits
  - Shift right by  $i$  bits = divides by  $2^i$  (unsigned only)  
(much faster than division, used in compiler optimization)
  - Example: **srl x9, x20, x21**

0000000	10101	10100	101	01001	0110011
---------	-------	-------	-----	-------	---------

- sra: shift right arithmetic (R-type)
  - Shift rs1 right and fill with sign bits (no **s1a**!)
  - Shift right by  $i$  bits = divides by  $2^i$  (signed)
  - Example: **sra x9, x20, x21**

*Note the  
encoding*

0100000	10101	10100	101	01001	0110011
---------	-------	-------	-----	-------	---------

- **x20** = 1111 ... 1000; **x21** = 0000 ... 0011 = 3
- After instruction executed, **x9** will contain **1111** ... 1111



# Shift Immediate Operations

- slli: shift left logical immediate (32-bit RISC-V: RV32)

0000000	shamt	rs1	001	rd	0010011
7 bits	5 bits	5 bits	3 bits	5 bits	7 bits

- *shamt*: how many bits to shift
  - 32-bit data can only be shifted 32 bit positions → 5-bit shamt
  - Same format as R-type, except rs2 becomes shamt
- slli: how about 64-bit RISC-V (RV64)?
  - We can shift 64 bit positions and need 6 bits for shamt
  - How?

000000	shamt	rs1	001	rd	0010011
6 bits	6 bits	5 bits	3 bits	5 bits	7 bits



# Logic Operations

- **and**: useful for setting certain portions of a word to 0s, while leaving the rest alone → *masking*

– Example: **and x9, x10, x11**

x10	00000000 00000000 00000000 00000000 00000000 00000000 00001101 11000000
x11	00000000 00000000 00000000 00000000 00000000 00000000 00111100 00000000
x9	00000000 00000000 00000000 00000000 00000000 00000000 00001100 00000000

- **or**: useful for forcing certain bits of a word to 1s, while leaving others unchanged

– Example: **or x9, x10, x11**

x10	00000000 00000000 00000000 00000000 00000000 00000000 00001101 11000000
x11	00000000 00000000 00000000 00000000 00000000 00000000 00111100 00000000
x9	00000000 00000000 00000000 00000000 00000000 00000000 00111101 11000000



# Logic Operations

- **xor**: differencing operation
  - Set some bits to 1, leave others unchanged

**xor x9, x10, x12**

XOR		
0	0	0
0	1	1
1	0	1
1	1	0

x10	00000000 00010100 00000000 00000000 00000000 00000000 00001101 11000000
x12	11111111 11000011 11111111 11111111 11111111 11111111 11111111 11111111
x9	11111111 11010111 11111111 11111111 11111111 11111111 11110010 00111111

- If mask (x12) = 1,  $0 \rightarrow 1$  and  $1 \rightarrow 0$
- If mask (x12) = 0,  $0 \rightarrow 0$  and  $1 \rightarrow 1$
- So, if mask is all 1s, XOR becomes a NOT operation  
→ no need for a **not** instruction (2-address) in RISC-V





# So Far...

- All instructions studied so far have allowed us to manipulate data
- So we have built a calculator
- However, in order to build a computer, we need the ability to make decisions and change the execution flow



# Conditional Branch Operations

- Branch to a *labeled instruction* if a condition is true
  - Otherwise, continue sequentially
- **beq rs1,rs2,L1** ← Memory address
  - if (rs1 == rs2) branch to instruction labeled L1, i.e., load memory address of that instruction to **PC** (Program Counter), so that the CPU will next fetch and execute the branch-to instruction
- **bne rs1,rs2,L1**
  - If (rs1 != rs2) branch to instruction labeled L1 by storing L1 (an address) into PC



# Compiling If Statements

- C code:

```
if (i==j) f = g + h;  
else f = g - h;
```

- f, g, h, i, j in  
x19, x20, x21, x22, x23

- Compiled RISC-V code:


```
bne x22,x23,Else
```

```
add x19,x20,x21
```

```
beq x0,x0,Exit // unconditional
```

```
Else: sub x19,x20,x21
```

```
Exit: ...
```



*Why not unconditional jump?*

Assembler calculates address  
(discussed later)



# Compiling Loop Statements

- C code:

```
while (save[i] == k) i += 1;
```

– **i** in **x22**, **k** in **x24**, base address of **save**[] in **x25**

- Compiled RISC-V code:

```
Loop:  slli  x10,x22,3    // x10←i*8  
       add   x10,x10,x25 // x10←save[i]  
       ld    x9,0(x10)  
       bne   x9,x24,Exit  
       addi  x22,x22,1  
       beq   x0,x0,Loop
```


```
Exit:  ...
```





# Concept of Basic Blocks

- A **basic block** is a sequence of instructions with
  - No embedded branches (except at end)
  - No branch targets (except at beginning)
- A compiler identifies basic blocks for optimization
  - All instructions in a basic block will surely be executed once the basic block is entered
    - can rearrange execution seq.
- An advanced processor can accelerate execution of basic blocks with hardware scheduling



# More Conditional Branch Operations

- **blt rs1,rs2,L1**
  - if ( $rs1 < rs2$ ) branch to instruction labeled L1
- **bge rs1,rs2,L1**
  - if ( $rs1 \geq rs2$ ) branch to instruction labeled L1
- Example:
  - if (a < b) a += 1;**
  - a in **x22**, b in **x23**

```
bge    x23,x22,Exit    // branch if b >= a
addi   x22,x22,1
```

**Exit:**



# Signed vs. Unsigned Comparison

- Signed comparison: **blt, bge**
- Unsigned comparison: **bltu, bgeu**
- Example:
  - $x22 = 1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1111$
  - $x23 = 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0001$
  - **blt x22,x23,L1 # signed**  
     $-1 < +1 \Rightarrow$  branch to L1
  - **bltu x22,x23,L1 # unsigned**  
     $+4,294,967,295 > +1 \Rightarrow$  do not branch





# Branch Instruction Design

- Alternative 1: (MIPS approach)
  - Set a register based on comparison result (e.g., **slt**), then branch on the value in that register with **beq** or **bne**
  - Pro: hardware for  $<$ ,  $\geq$ , ... slower than  $=$ ,  $\neq$ , because combining with branch involves more work per instruction, requiring a slower clock  $\rightarrow$  all instructions penalized!
  - Con: need two instructions instead of one
- Alternative 2: (ARM approach)
  - Keep extra bits (*condition codes* or *flags*) that record what occurred during an instruction, e.g., overflow, negative, ...
  - Conditional branches then test on condition codes
  - Con: will create dependencies on condition codes

