

EECS4030: Computer Architecture

Arithmetic for Computers (I)

Prof. Chung-Ta King

Department of Computer Science

National Tsing Hua University, Taiwan

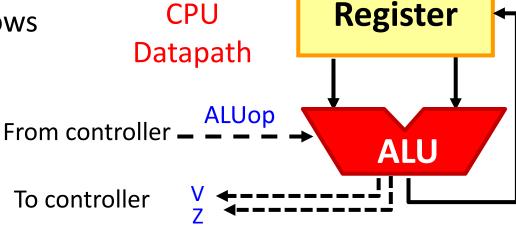
(Adapted from textbook slides https://www.elsevier.com/books-and-journals/book-companion/9780128122754/lecture-slides)

Outline

- Addition and subtraction (Sec. 3.2)
 - Constructing a basic ALU (Sec. A.5)
- Multiplication (Sec. 3.3)
- Division (Sec. 3.4)
- Floating point (Sec. 3.5)
- Parallelism and computer arithmetic: subword parallelism (Sec. 3.6)
- Streaming SIMD extensions and advanced vector extensions in x86 (Sec. 3.7)
- Subword parallelism and matrix multiply (Sec. 3.8)

Constructing a Basic ALU

- Support the following arithmetic and logic operations
 - add, sub: two's complement addition/subtraction with overflow detection
 - and, or, nor: bitwise logical AND, logical OR, logical NOR
 - slt (set-less-than): compare two input numbers (by subtraction) and output a 1 if first is less than second, otherwise output a 0
 - V: result overflows
 - Z: result is zero



Basic ALU and RISC-V Instructions

- ALU operations for executing RISC-V instructions
 - slt (set-less-than): slt, sltu, slti, sltiu, ...

 slt x9,x21,x22 $x9 \leftarrow 1$, if x21 < x22 $x9 \leftarrow 0$, else
 - add, sub: add, addi, sub, beq, bne, blt, bge, ld, sd, slt, sltu, slti, sltiu,...
 - and, or: and, andi, or, ori, ...
 - nor: not used in RISC-V
 - $\frac{7}{\text{beq}}$...

 beq $\frac{x21}{x22}$, $\frac{x22}{L1}$ jump to $\frac{1}{L1}$ if $\frac{x21}{L1} = \frac{x22}{L1}$ check if $\frac{x21}{L1} = \frac{x22}{L1}$

Functional Specification

ALU Control (ALUop)	Function	
0000	and	
0001	or	
0010	add	
0110	sub	
0111	set-less-than	
1100	nor	

Fig. A.5.13

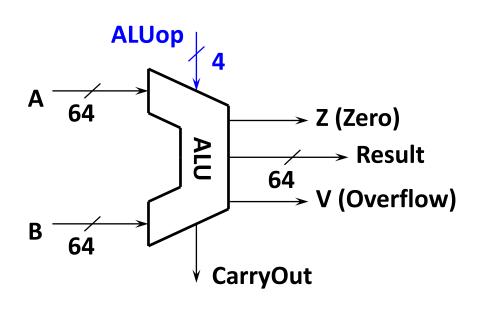
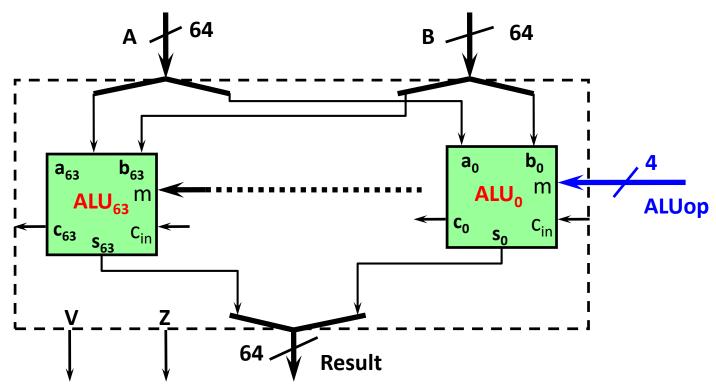


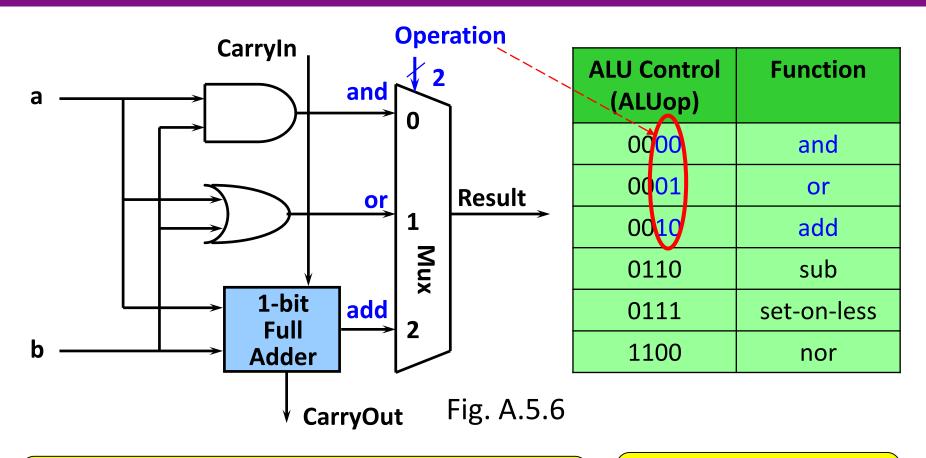
Fig. A.5.14

Design Strategy: Bit-slice ALU

- Design trick: divide and conquer
 - Break the problem into simpler problems, solve individual problems and glue together the solutions



Zoom-in to 1-bit ALU



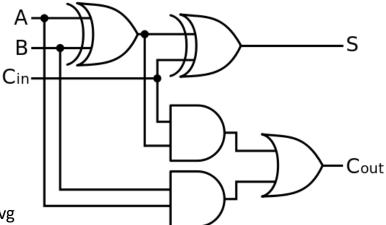
All logic gates operate to produce results and only the desired one is chosen

Only combinational logic

Recall 1-Bit Full Adder

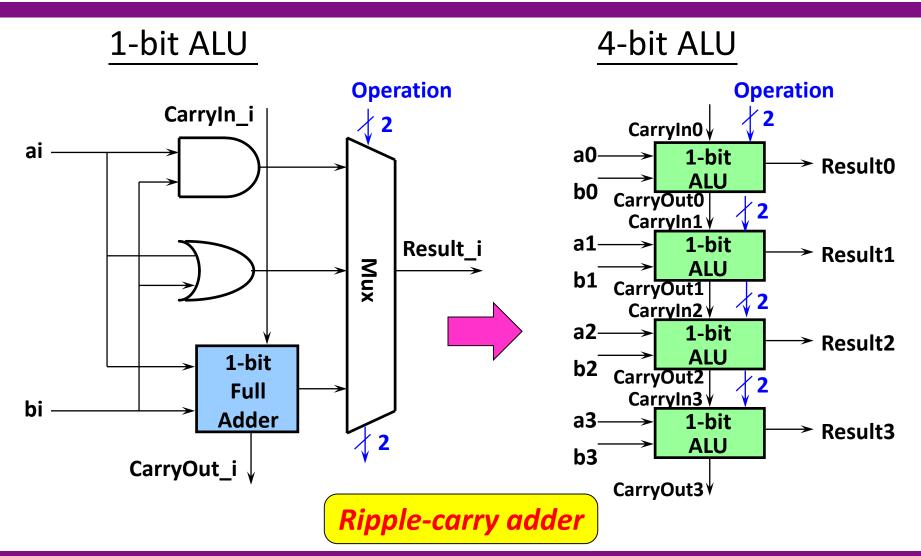
Inputs		Outputs		
а	b	Carryin	CarryOut	Sum
0	0	0	0	0
0	0	1	0	1
0	1	0	0	1
0	1	1	1	0
1	0	0	0	1
1	0	1	1	0
1	1	0	1	0
1	1	1	1	1

Fig. A.5.3



https://commons.wikimedia.org/wiki/File:Full-adder.svg

From 1-bit to 4-bit ALU

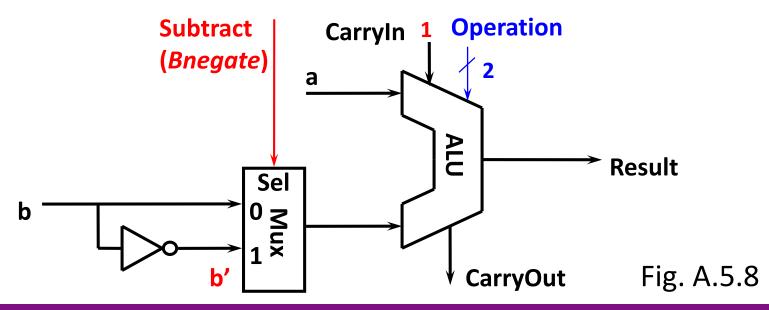


How about Subtraction?

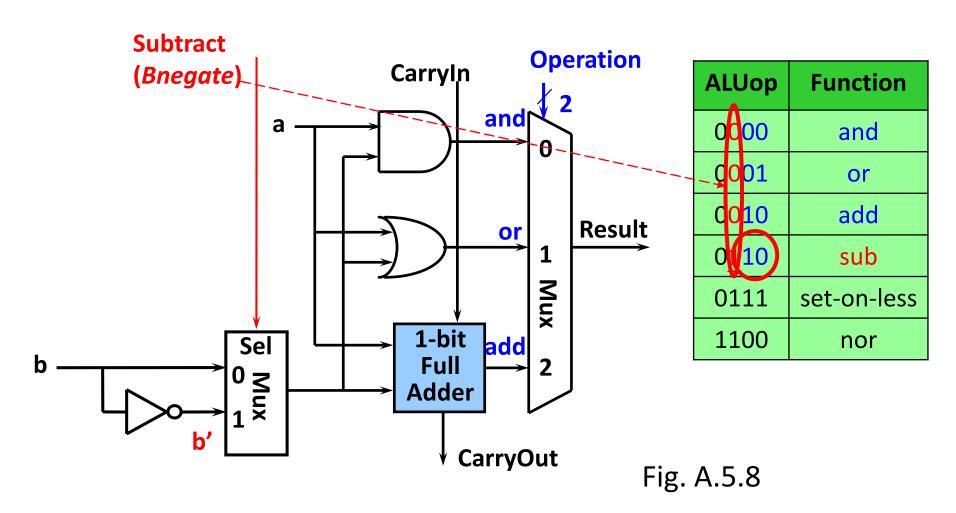
- 2's complement operation: take inverse of every bit and add 1 (at c_{in} of first stage)
- So, subtraction becomes addition of 2's complement

$$-A-B=A+(-B)=A+(B'+1)=A+B'+1$$

B': bit-wise inverse of B

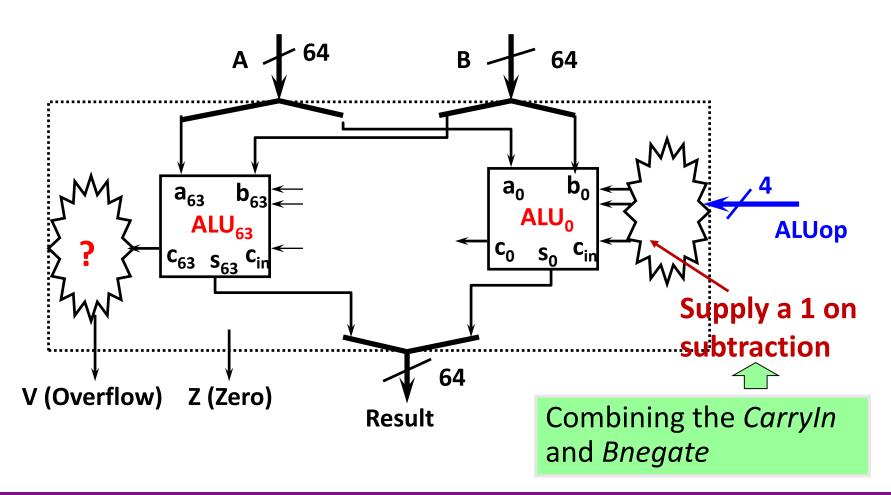


Zoom-in to 1-bit ALU

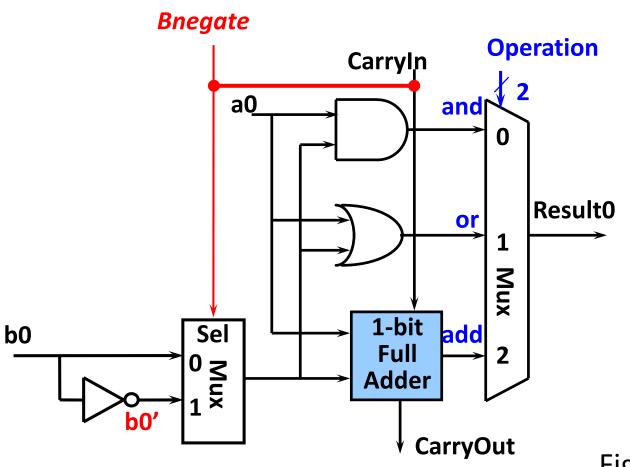


Revised Diagram

LSB and MSB need to do a little extra



Zoom-in to Rightmost Bit of ALU



ALUop	Function	
0000	and	
0001	or	
0010	add	
0 <u>1</u> 10	sub	
0111	set-on-less	
1100	nor	

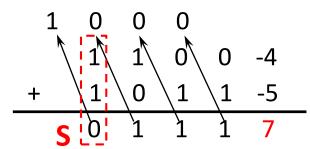
Fig. A.5.8

Overflow of Signed Integer Addition

Decimal	Binary	Decimal	2's complement
0	0000	0	0000
1	0001	-1	1111
2	0010	-2	1110
3	0011	-3	1101
4	0100	-4	1100
5	0101	-5	1011
6	0110	-6	1010
7	0111	-7	1001
		-8	1000

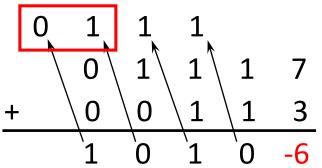
Ex:
$$7 + 3 = 10$$
 but ...

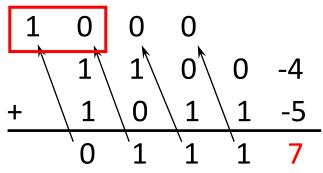
$$-4 - 5 = -9$$
 but ...



Overflow Detection

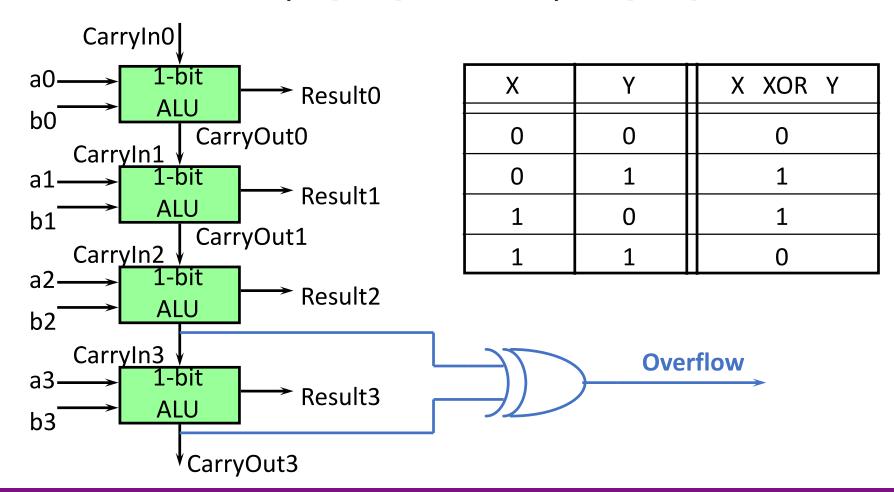
- Overflow: result too big/small to represent
 - -8 ≤ 4-bit binary number ≤ 7
 - Adding operands with different signs → no overflow
 - Overflow occurs when adding:
 - 2 positive numbers and the sum is negative
 - 2 negative numbers and the sum is positive
 - → sign bit is set with the value of the result
 - Overflow: if Carry into MSB ≠ Carry out of MSB



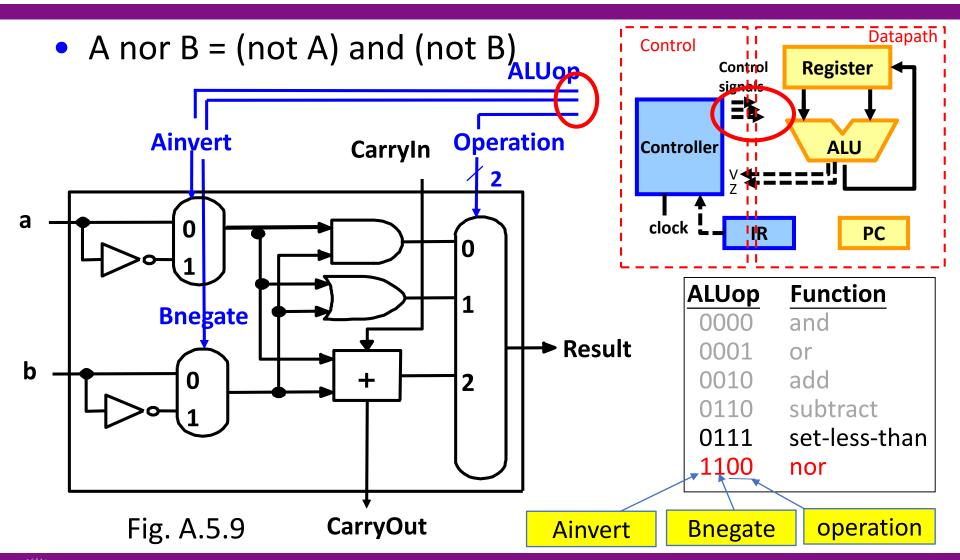


Overflow Detection Logic

Overflow = CarryIn[N-1] XOR CarryOut[N-1]

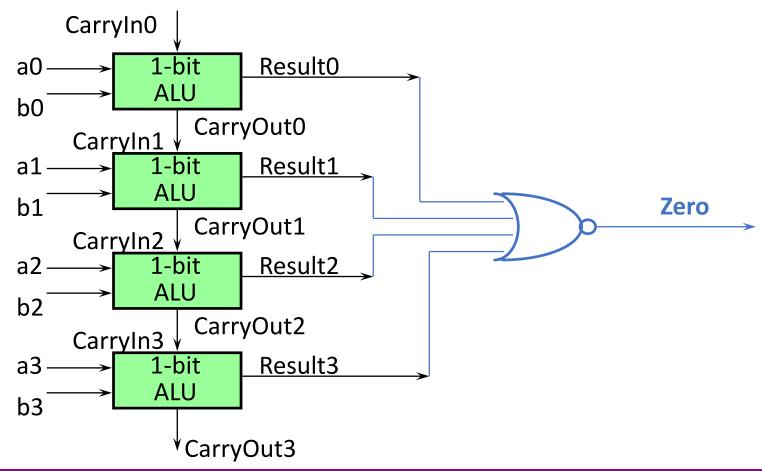


NOR Operation



Zero Detection Logic

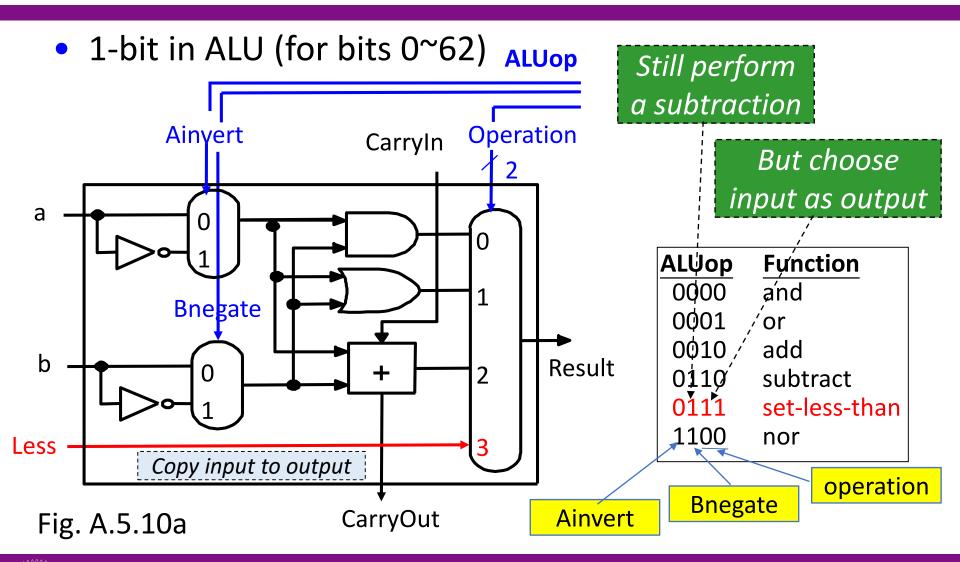
To support conditional jump: by a one BIG NOR gate



Set-Less-Than

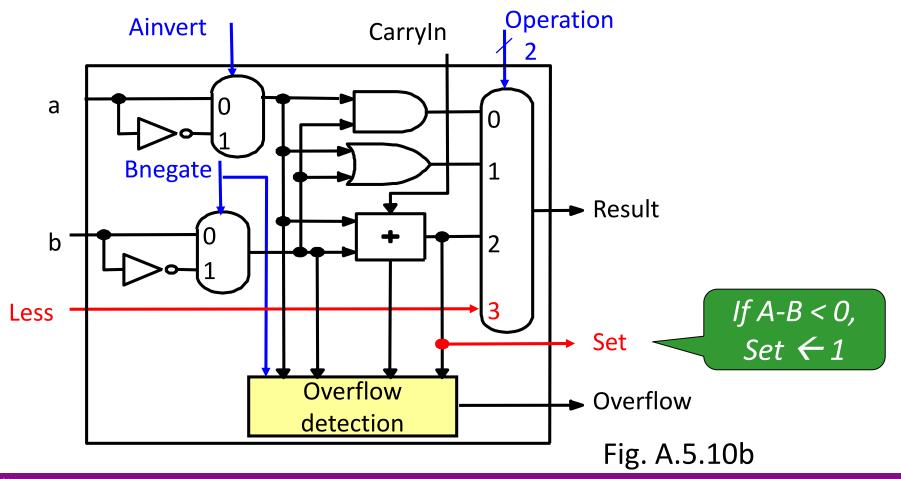
- If A < B, then Result = 00..01, else Result = 00..00
- If (A B) < 0, then ...
- If A − B & sign of Result is 1, then ...
- If A B & Result₆₃ = 1, then ...
- If A B & Result₆₃ = 1, then Result₀ = 1, else Result₀ = 0, and Result₆₃ \sim Result₁ = 0
- $A B \& Result_0 = Result_{63} \& Result_{63} \sim Result_1 = 0$
 - No need to do if-then-else test

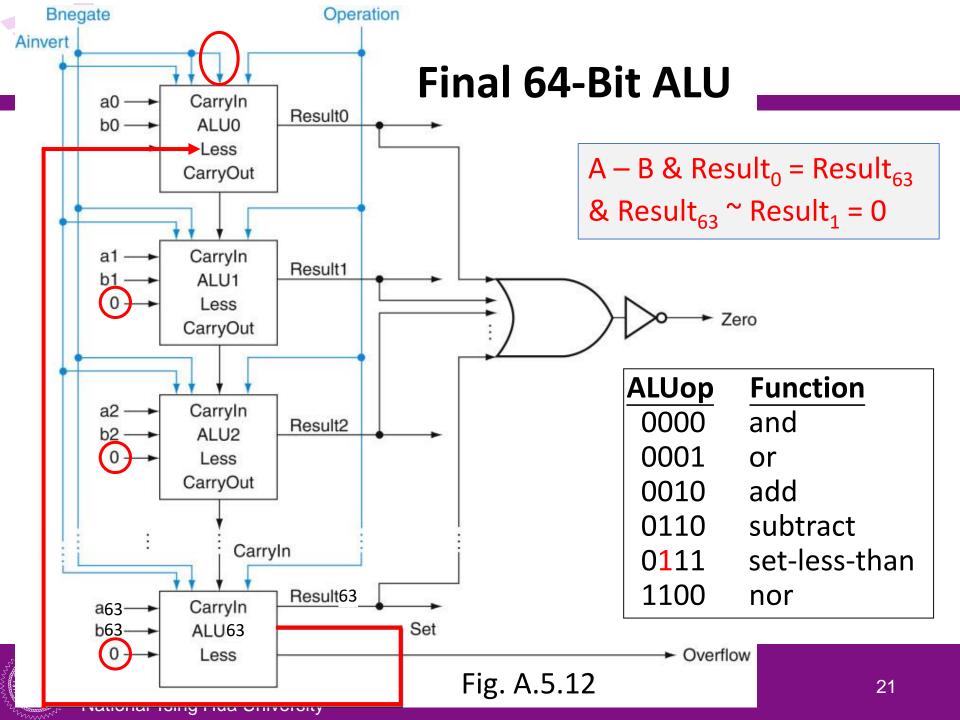
Set-Less-Than (I)

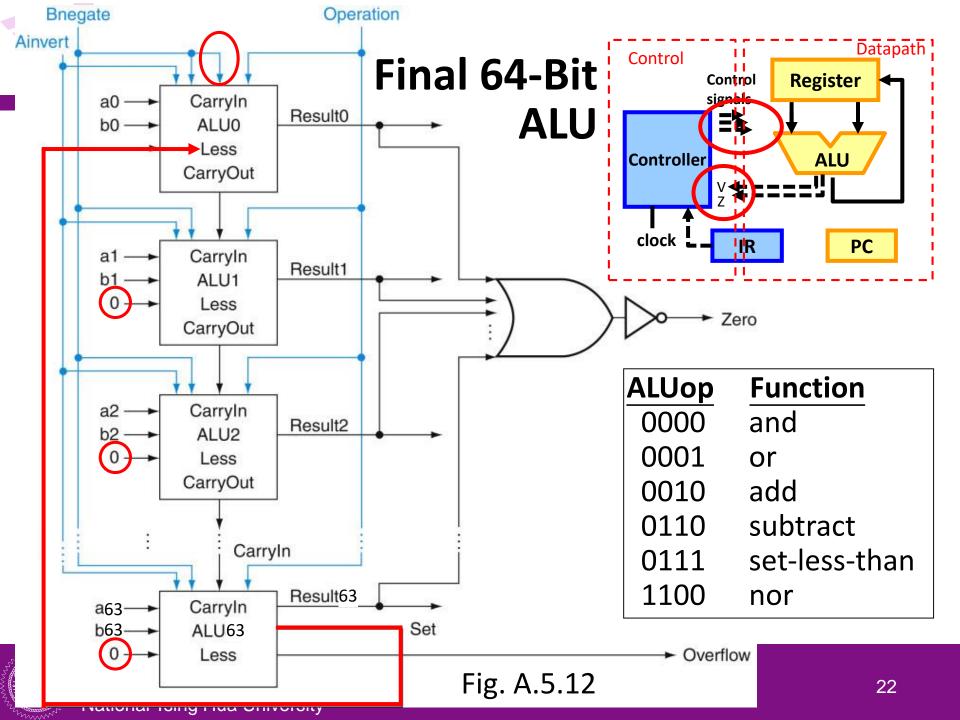


Set-Less-than (II)

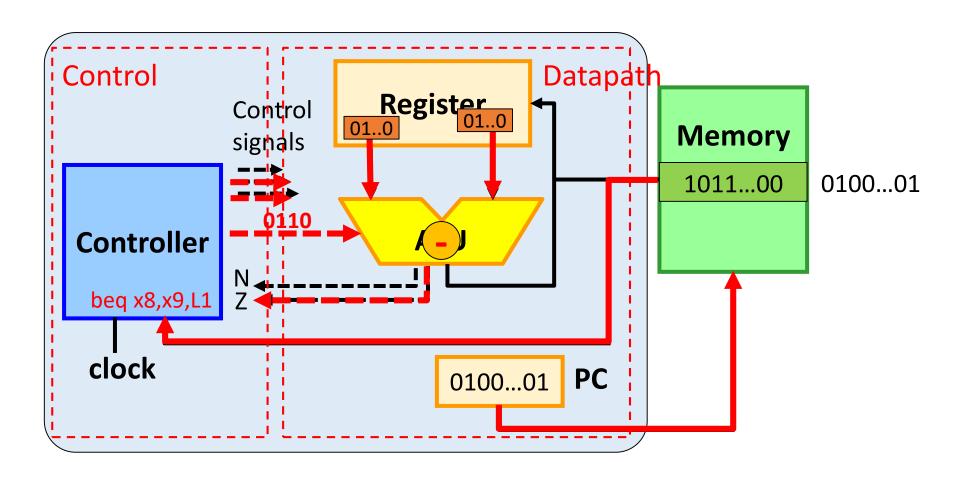
Sign bit in ALU (bit 63)







Putting It Altogether



Verilog Behavioral Definition of 64-Bit ALU

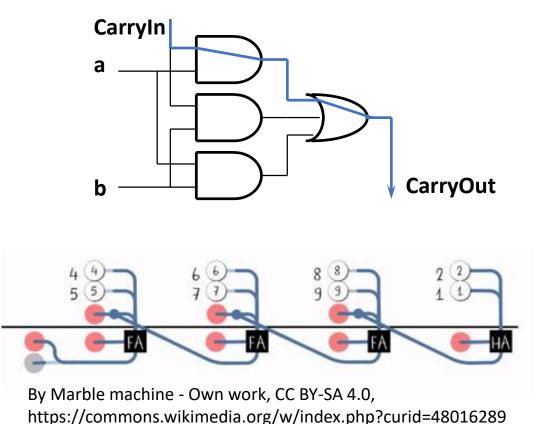
```
module RISCVALU(ALUctl, A, B, ALUOut, Zero);
  input [3:0] ALUctl;
  input [63:0] A,B;
  output reg [63:0] ALUOut;
  output Zero;
  assign Zero = (ALUOut==0); //Zero true if ALUOut 0
  always @(ALUctl, A, B) begin //reevaluate if change
    case (ALUctl)
      0: ALUOut <= A & B;
                                          ALUop
                                                 Function
      1: ALUOut <= A | B;
                                                 and
                                           0000
      2: ALUOut <= A + B;
      6: ALUOut <= A - B;
                                           0001
                                                 or
                                           0010
                                                 add
      7: ALUOut (= A < B ? 1 : 0;
                                           0110
                                                 subtract
      12: ALUOut <= ~(A | B); // nor
                                           0111 set-less-than
      default: ALUOut <= 0;</pre>
                                           1100
                                                 nor
    endcase
  end
```

Fig. A.5.15

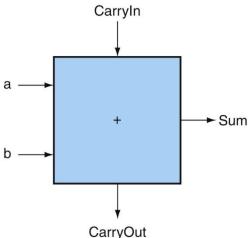
endmodule

Problems with Ripple Carry Adder

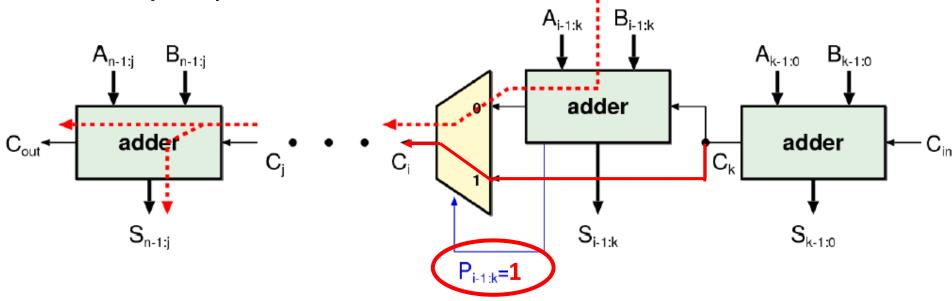
- Carry bit may have to propagate from LSB to MSB
 → worst case delay: N-stage delay
- CarryIn0 a0 1-bit Result0 b0 CarryOut0 CarryIn1 a1 1-bit Result1 b1 CarryOut1 CarryIn2 a2 1-bit Result2 b2 CarryOut2 CarryIn3 a3 1-bit Result3 **b**3 CarryOut3



- Strategy: determine carry in to high-order bits sooner
- How?
 - Examine how carry is generated and propagated
 - Carry generation: under what condition a carry is generated at a bit stage?
 - If a=1 and b=1, a carry is generated no matter what the value of carry input is Generate = g = a & b
 - Carry propagation: under what condition a carry can be propagated at a bit stage?
 - If only a=1 or only b=1, the value of carry input is propagated to the carry output
 Propagate = p = a xor b



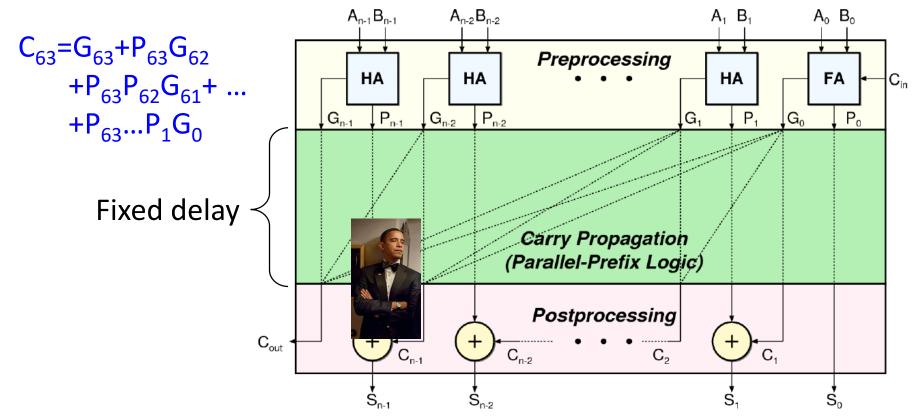
Carry skip adder:



- $-P_{i-1:k} = P_{i-1} \& P_{i-2} \& ... \& P_k = 1 \rightarrow C_i = C_k$
- $P_{i-1:k}$ = 0 → this stage will not propagate carry → the result of the carry is generated within this stage

http://www.syssec.ethz.ch/education/Digitaltechnik_14

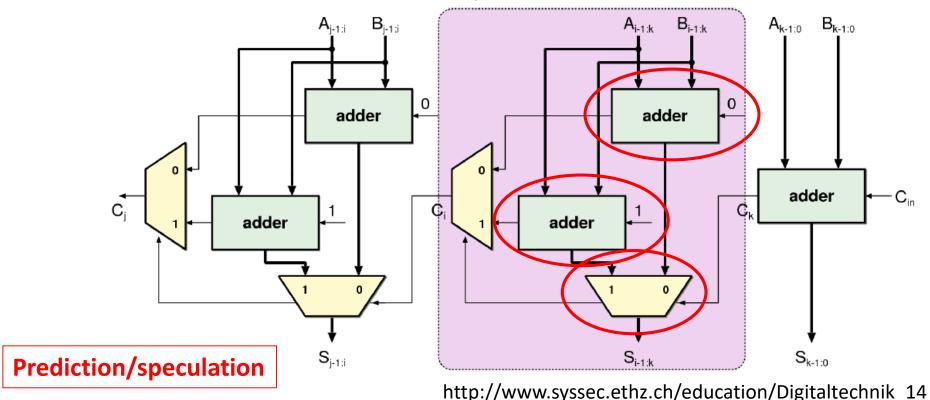
- Carry lookahead adder:
 - Look ahead and anticipate the carry to come (Sec. B.6)



http://www.syssec.ethz.ch/education/Digitaltechnik_14



- Carry select adder:
 - Design principle: parallelism (do both paths first and choose the desired result after carry in is known)



Outline

- Addition and subtraction (Sec. 3.2)
- Multiplication (Sec. 3.3)
- Division (Sec. 3.4)
- Floating point (Sec. 3.5)
- Parallelism and computer arithmetic: subword parallelism (Sec. 3.6)
- Streaming SIMD extensions and advanced vector extensions in x86 (Sec. 3.7)
- Subword parallelism and matrix multiply (Sec. 3.8)

Unsigned Multiply

Paper and pencil example (unsigned):

```
1000 Multiplicand

X 1001 Multiplier

00001000
00000000
01000000

01001000 Product
```

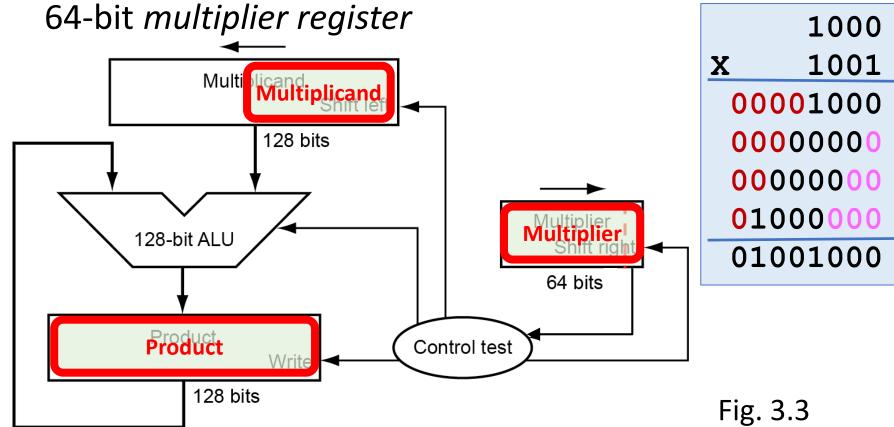
Binary makes it easy:

```
- 0 → place 0 (0 x multiplicand)
- 1 → place a copy (1 x multiplicand)
```

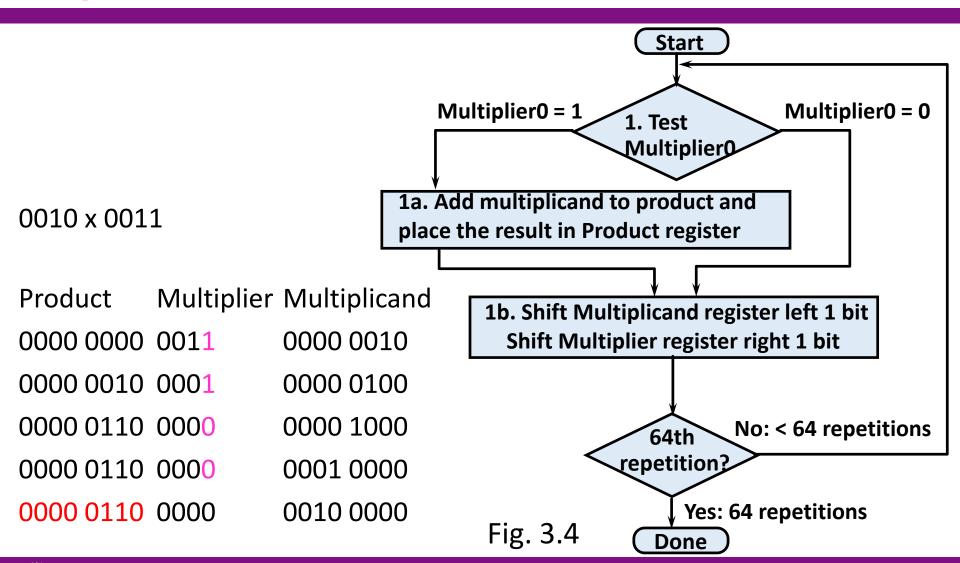
m bits x n bits = m+n bit product

Unsigned Multiplier for RISC-V (Ver. 1)

 128-bit multiplicand register (with 64-bit multiplicand at right half), 128-bit ALU, 128-bit product register,



Multiply Algorithm (Ver. 1)

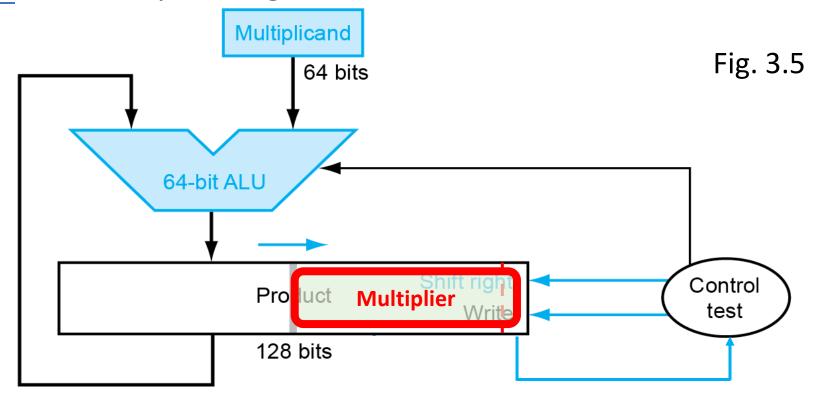


Observations: Multiply Ver. 1

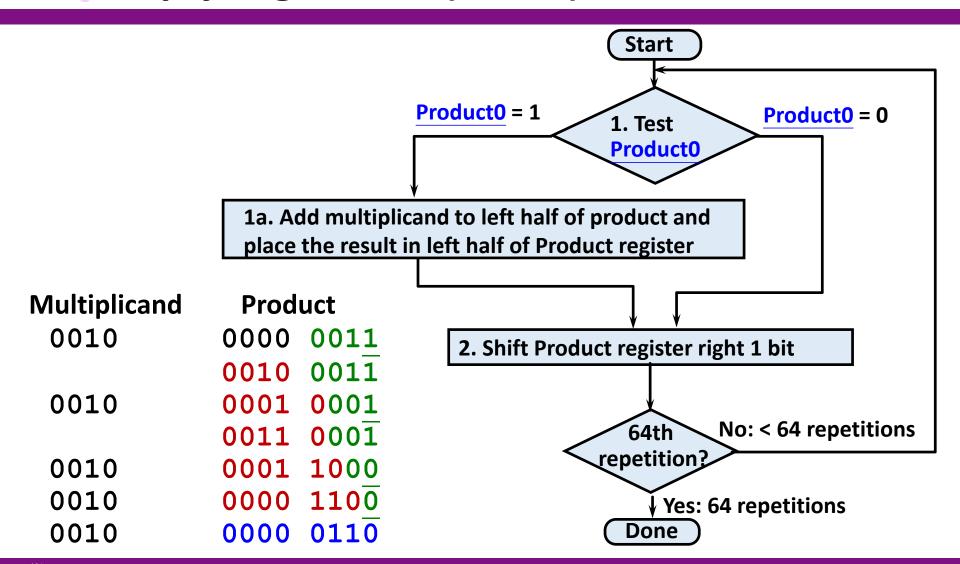
- 1 clock per iteration → ~64 clocks per multiply
 - Ratio of multiply to addition 5:1 to 100:1
- Half of the bits in multiplicand always 0
 - → 128-bit adder is wasteful
- 0's inserted in right of multiplicand as shifted
 - → least significant bits of product never changed once formed
- Instead of shifting multiplicand to left, shift product to right?

Unsigned Multiplier for RISC-V (Ver. 2)

64-bit Multiplicand register, 64-bit ALU, 128-bit
 Product register (with 64-bit multiplier at right half),
 0-bit Multiplier register



Multiply Algorithm (Ver. 2)



Observations: Multiply Ver. 2

- 2 steps in ver. 2, because Product register needs to be written twice (step 1a and step 2)
 - Steps 1a and 2 can be combined to complete in 1 cycle through hardware wiring
- What about signed multiplication?
 - The easiest solution is to make both positive and remember whether to complement product when done (leave out sign bit, run for 63 steps)
 - Apply definition of 2's complement
 - sign-extend partial products and subtract at end
 - Booth's algorithm: use same HW and save cycles

Faster Multiplier

1000 x 1111

Use multiple adders: cost/performance tradeoff

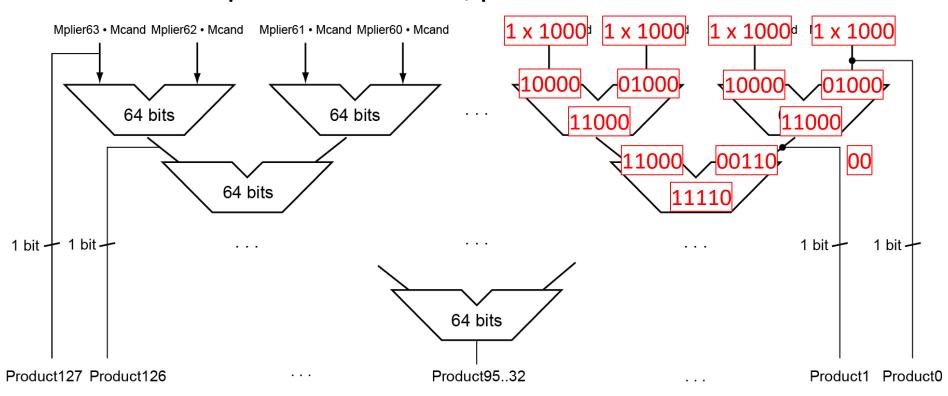


Fig. 3.7

Can be pipelined: several multiplications in parallel

RISC-V Integer Multiplica To get 128-bit product:

To get 128-bit product:

mulh[[s]u] rdh, rs1, rs2

mul rdl, rs1,rs2

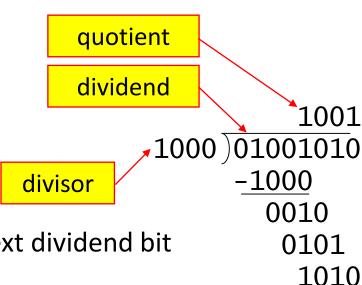
- mul: multiply
 - Gives lower 64 bits of the product in destination register
- mulh: multiply high
 - Gives upper 64 bits of the product in destination register, assuming signed × signed
- mulhu: multiply high unsigned
 - Gives upper 64 bits of the product in destination register, assuming unsigned × unsigned
- mulhsu: multiply high signed/unsigned
 - Gives upper 64 bits of the product in destination register, assuming signed × unsigned
- Use mulh result to check for 64-bit overflow

Outline

- Addition and subtraction (Sec. 3.2)
- Multiplication (Sec. 3.3)
- Division (Sec. 3.4)
- Floating point (Sec. 3.5)
- Parallelism and computer arithmetic: subword parallelism (Sec. 3.6)
- Streaming SIMD extensions and advanced vector extensions in x86 (Sec. 3.7)
- Subword parallelism and matrix multiply (Sec. 3.8)

Division

- Check for 0 divisor
- Long division approach
 - If divisor ≤ dividend
 - 1 bit in quotient, subtract
 - Otherwise
 - 0 bit in quotient, bring down next dividend bit
- Restoring division approach
 - Do the subtract, and if remainder goes < 0, add divisor back
- Signed division
 - Divide using absolute values
 - Adjust sign of quotient and remainder as required



remainder

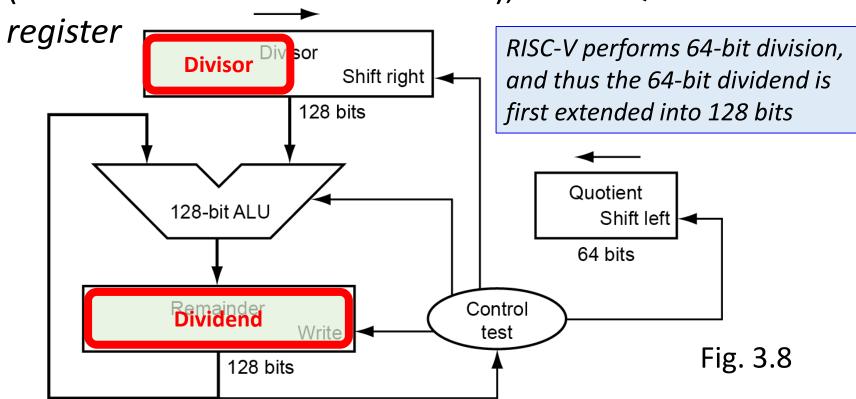
n-bit operands yield *n*-bit quotient and remainder

-1000

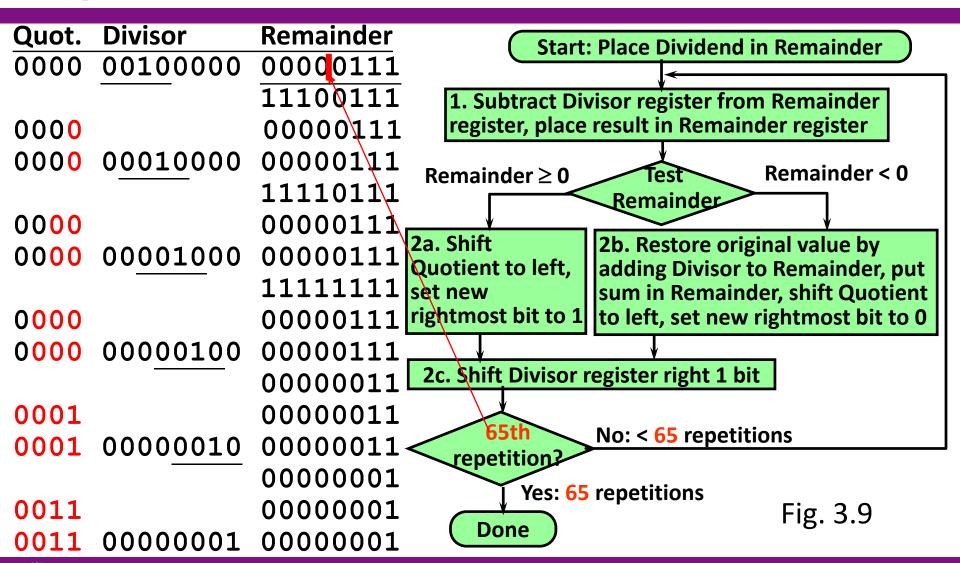
0010

Divide Hardware for RISC-V (Version 1)

 128-bit Divisor register (initialized with 64-bit divisor in left half), 128-bit ALU, 128-bit Remainder register (initialized with 128-bit dividend), 64-bit Quotient



Divide Algorithm (Version 1)



Observations: Divide Version 1

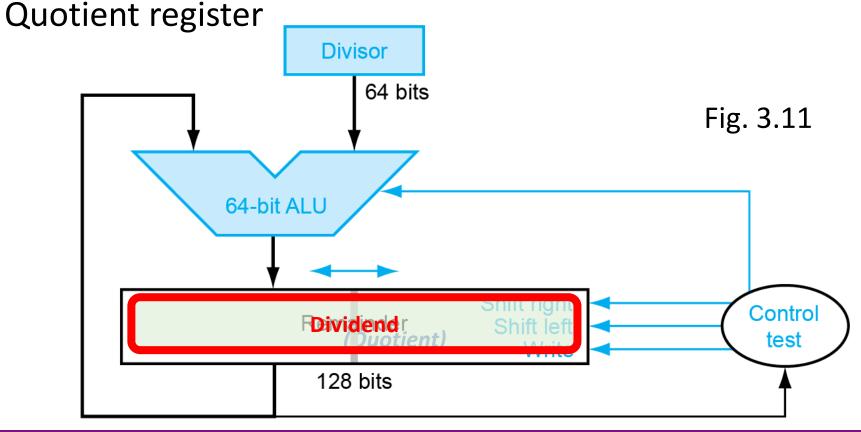
- Half of the bits in divisor register always 0
 - \rightarrow 1/2 of 128-bit adder is wasted
 - \rightarrow 1/2 of divisor is wasted
- Instead of shifting divisor to right, shift remainder to left?
- For 64-bit division, 1st step cannot produce a 1 in quotient bit (otherwise quotient is too big for the register)
 - > switch order to shift first and then subtract
 - → save 1 iteration
- Eliminate Quotient register by combining with Remainder register as shifted left

Remainder

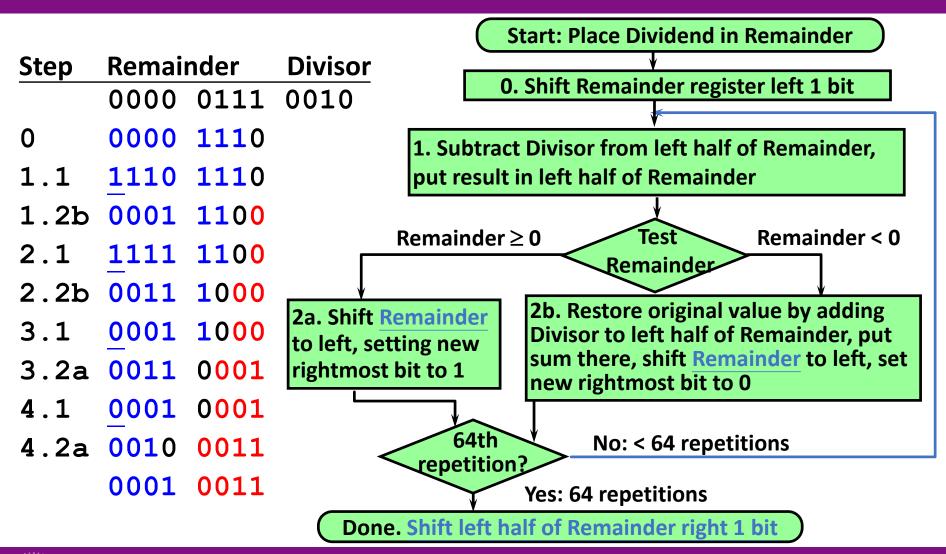
0000 0111

Divide Hardware (Version 2)

• 64-bit Divisor register, 64-bit ALU, 128-bit Remainder register (initialized with 128-bit dividend), <u>0</u>-bit



Divide Algorithm (Version 2)



Signed Divide

- Remember signs, make positive, complement quotient and remainder if necessary
- Alternative:

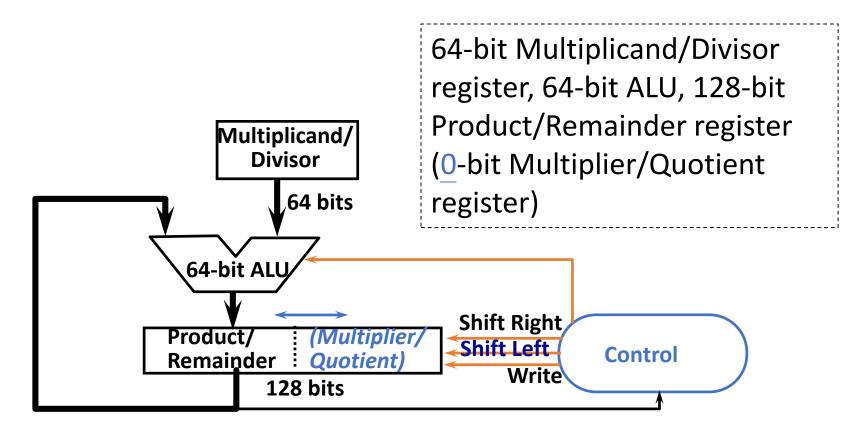
Let Dividend and Remainder have same sign, negate Quotient if Divisor sign and Dividend sign disagree

- e.g.,
$$-7 \div 2 = -3$$
, remainder = -1
-7 \div -2 = 3, remainder = -1

- Satisfy Dividend = Quotient x Divisor + Remainder
- Possible for quotient to be too large:
 - If divide 64-bit integer by 1, quotient is 64 bits

Observation: Multiply and Divide

 Same hardware as multiply: just need ALU to add or subtract, and 128-bit register to shift left or shift right



Faster Division

- Can't use parallel hardware as in multiplier
 - Subtraction is conditional on sign of remainder
- Faster dividers (e.g., SRT division) generate multiple quotient bits per step
 - Still require multiple steps

RISC-V Integer Division in M Extension

- div/divu: signed/unsigned division
 - Gives 64-bit signed/unsigned quotient in destination register
- rem/remu: signed/unsigned remainder
 - Gives 64-bit signed/unsigned remainder of the division
- If both quotient and remainder are required:

```
div[u] rdq, rs1, rs2
rem[u] rdr, rs1, rs2
```

- Microarchitectures can fuse them into a single division
- Overflow and division-by-zero don't produce errors
 - Just return defined results
 - Faster for the common case of no error