

EECS4030: Computer Architecture

The Processor (II)

Prof. Chung-Ta King

Department of Computer Science

National Tsing Hua University, Taiwan

(Adapted from textbook slides https://www.elsevier.com/books-and-journals/book-companion/9780128122754/lecture-slides)

Outline

- Problem statement and logic review (Sec. 4.1, 4.2)
- Building a simple RISC-V processor with datapath and control (Sec. 4.3, 4.4)
- Building a pipelined RISC-V processor with datapath and control (Sec. 4.5, 4.6)
- Dealing hazards in pipelined processor: data and control hazards (Sec. 4.7, 4.8)
- Handling exceptions (Sec. 4.9)
- More advanced topics: parallelism via instructions, ARM Cortex-A53 and Intel Core i7 Pipelines, instruction-level parallelism (Sec. 4.10, 4.11, 4.12)

Pipelining Is Natural!

Laundry example:

- Ann, Brian, Cathy, Dave each have one load of clothes to wash, dry, and fold
- Washer takes 30 minutes

- Dryer takes 40 minutes
- "Folder" takes 20 minutes

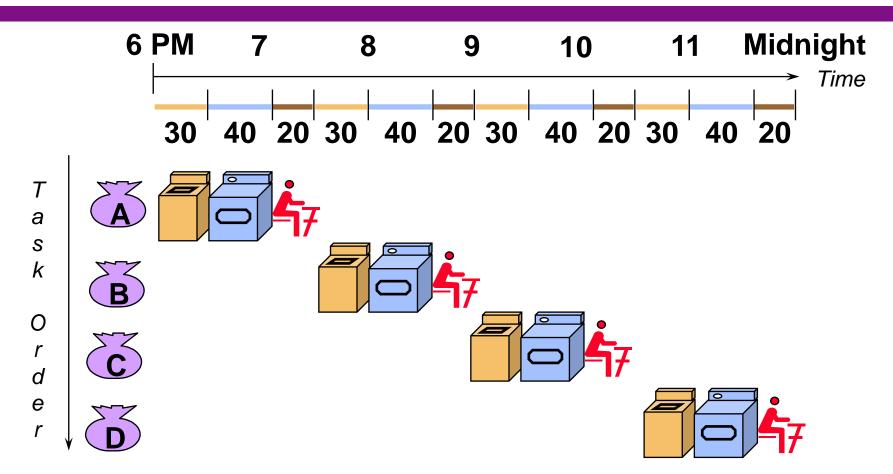






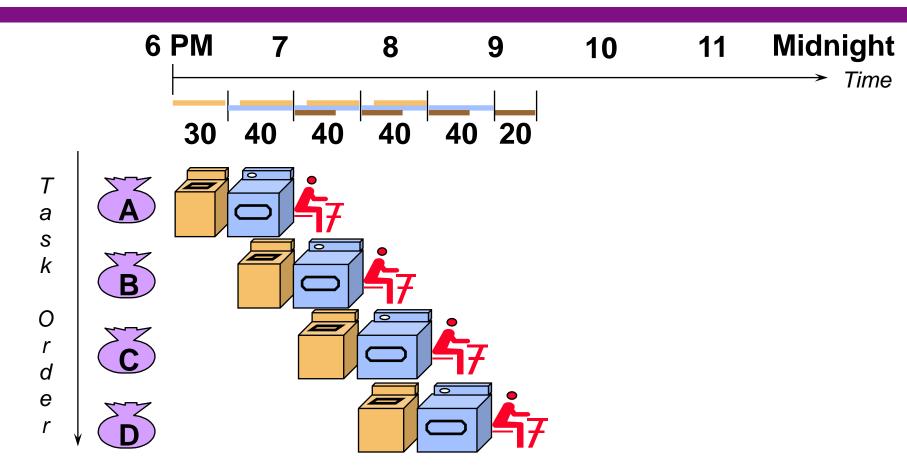


Sequential Laundry



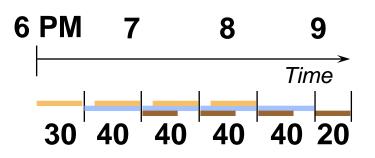
- Sequential laundry takes 6 hours for 4 loads
- If they learned pipelining, how long would it take?

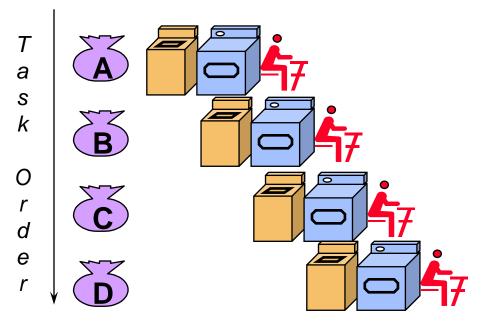
Pipelined Laundry: Start ASAP



- Pipelined laundry takes 3.5 hours for 4 loads
- Speedup = 8/3.5 = 2.3

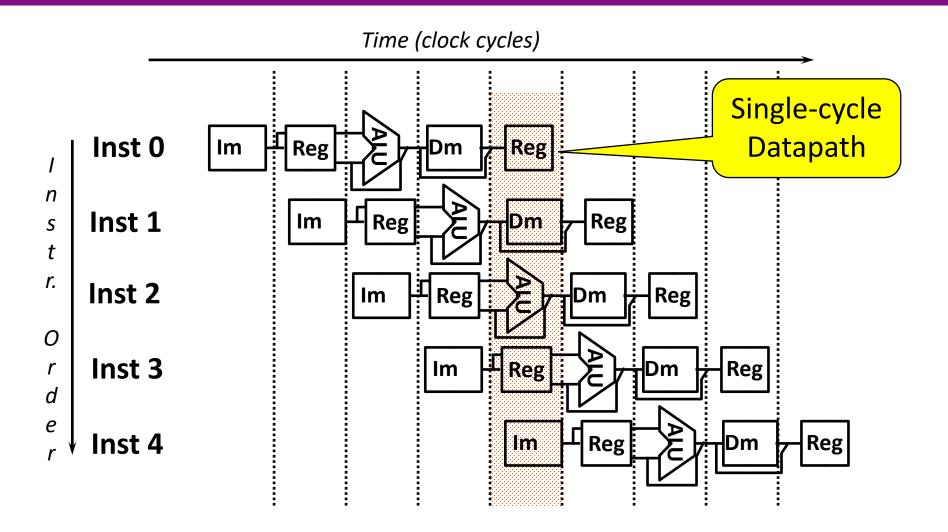
Pipelining Lessons





- Multiple tasks working at same time using different resources
- Doesn't help latency of single task, but throughput of entire
- Pipeline rate limited by slowest stage
 - → balance stage length
- Ideal speedup = # of stages
 - kN / (k+(N-1)): for k stages
 - Time to "fill" & "drain" the pipeline reduce speedup

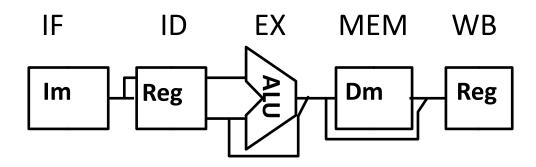
Why Pipeline? Because Resources Are There!



RISC-V Pipeline

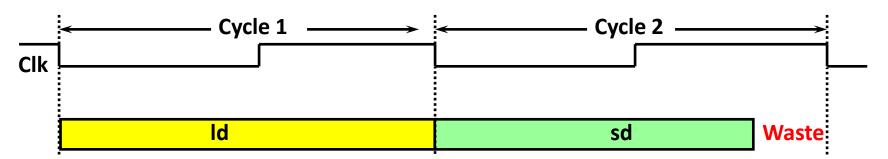
Five stages, one step per stage

- IF: Instruction fetch from memory
- ID: Instruction decode and register read
- EX: Execute operation or calculate address
- MEM: Access memory operand
- WB: Write result back to register

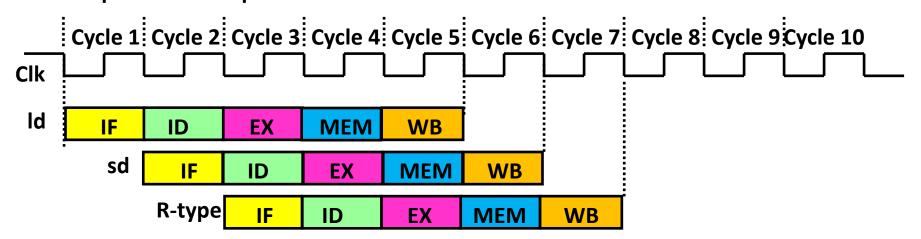


Single-Cycle vs. Pipeline

Single-cycle implementation:



Pipeline implementation:



Pipeline Performance

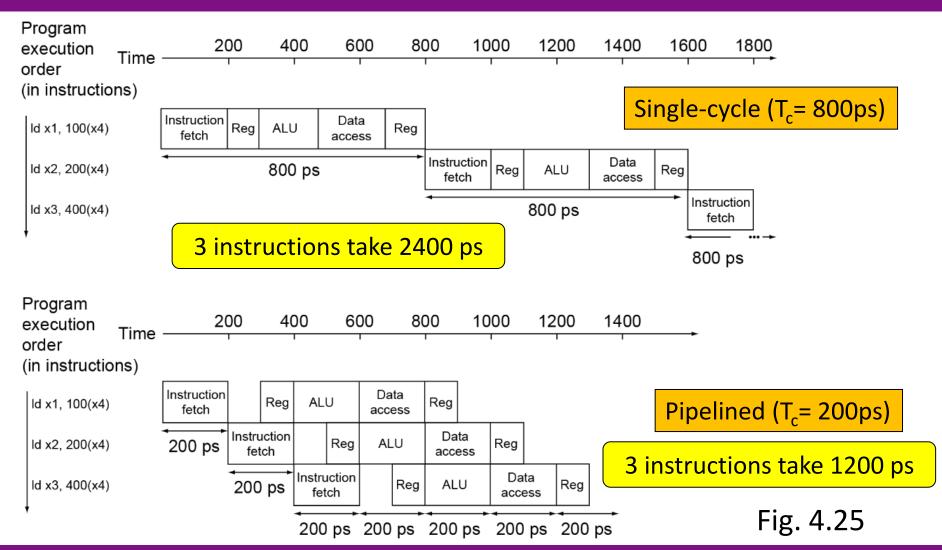
- Assume time for stages is
 - 100ps for register read or write
 - 200ps for other stages

Fig. 4.24

Instr	IF	ID	EX	MEM	WB	Total time
ld	200ps	100 ps	200ps	200ps	100 ps	800ps
sd	200ps	100 ps	200ps	200ps		700ps
R-type	200ps	100 ps	200ps		100 ps	600ps
beq	200ps	100 ps	200ps			500ps

- Compare pipelined with single-cycle datapath
 - Slowest stage vs. slowest instruction

Pipeline Performance



Pipeline Speedup

- If all stages are balanced,
 - i.e., all take the same time
 - Time between instructions_{pipelined}

= Time between instructions_{nonpipelined} =
$$800 \text{ ps}$$
 = 160 ps
Number of stages 5

Speedup for N ld instructions

= N * 800 ps
$$\cong$$
 5 (when N is large)
800 ps + (N-1)*160 ps

- If not balanced (as in previous slide), speedup is less
 - $(N * 800) / (800 + (N-1)*200) \approx 4$
- Speedup of pipeline is due to increased throughput
 - Latency (time for each instruction) does not decrease

Design a Pipelined RISC-V Processor

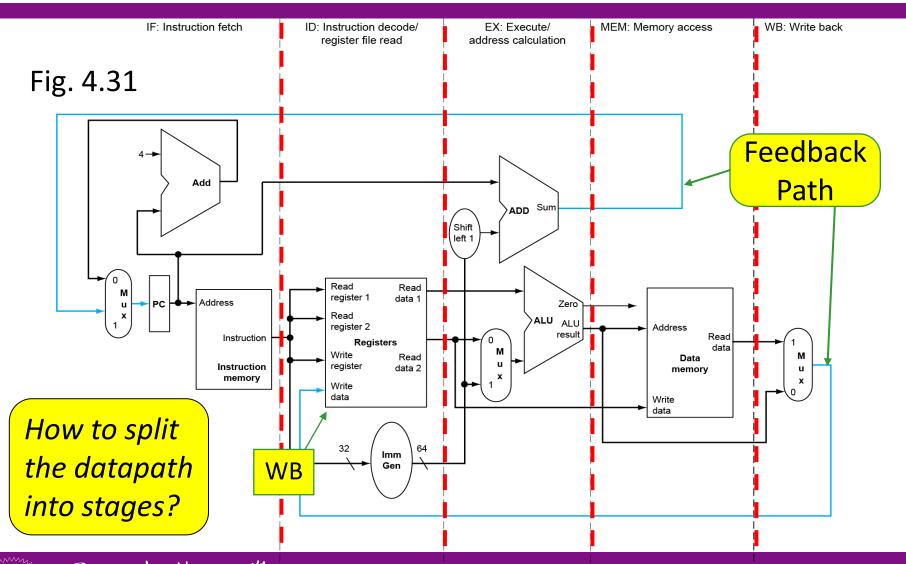
RISC-V ISA designed for pipelining

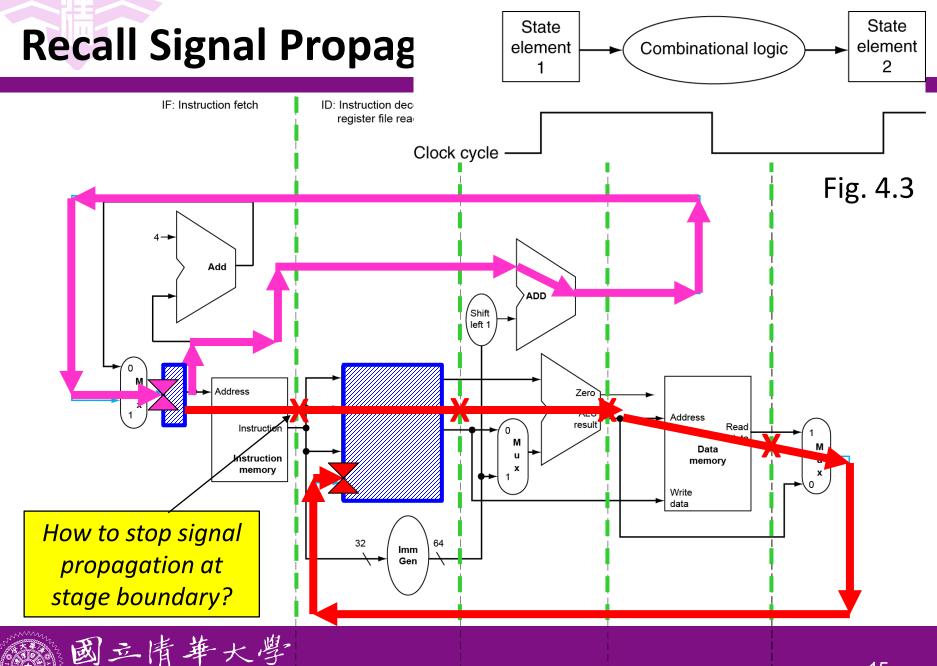
- All instructions are 32-bits
 - Easier to fetch and decode in one cycle
 - c.f. x86: 1- to 17-byte instructions
- Few and regular instruction formats
 - Can decode and read registers at the same step
- Load/store addressing
 - Can calculate address in the 3rd stage, access memory in the 4th stage

Design a Pipelined RISC-V Processor

- Examine the datapath and control diagram
 - Starting with single-cycle datapath
- Partition datapath into stages:
 - -5 stages: IF (instruction fetch), ID (instruction decode and register file read), EX (execution or address calculation), MEM (data memory access), WB (write back)
- Associate resources with stages
- Ensure that flows do not conflict, or figure out how to resolve
- Assert control in appropriate stage

Split Single-cycle Datapath



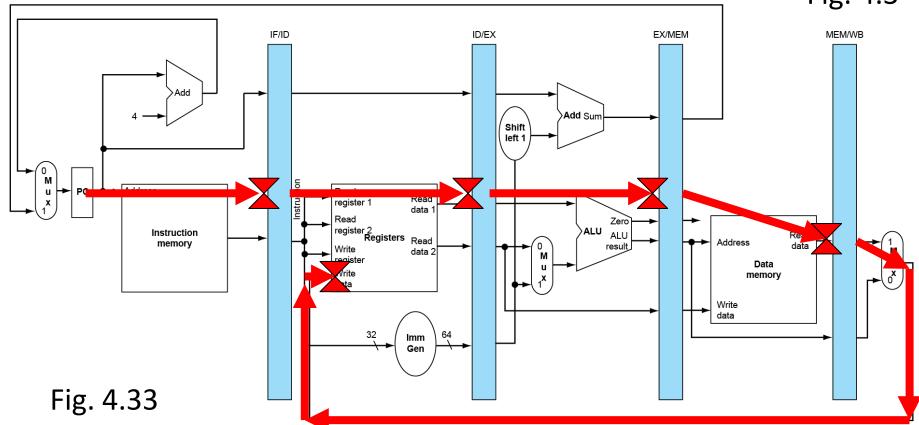


Idea: Add Registers

State element 1 State element 2

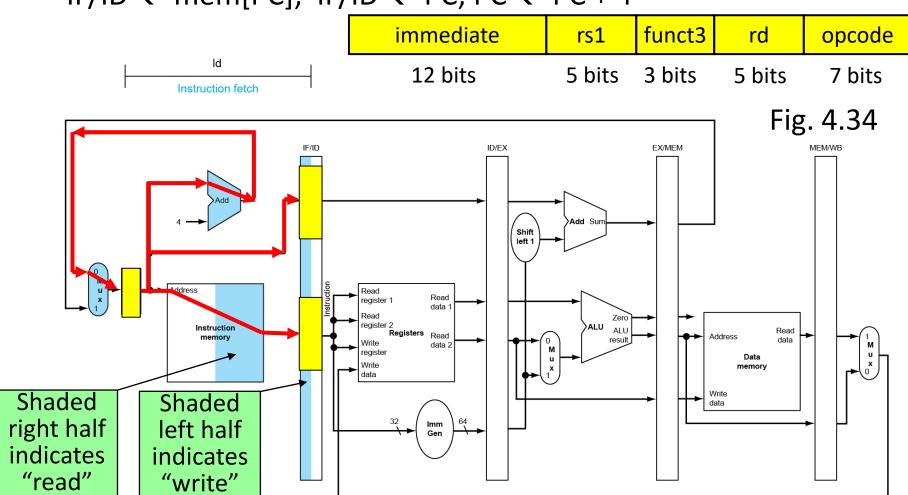
• Pipeline registers to ho "state" of an instruction "Clock cycle"

Fig. 4.3



IF Stage for Id

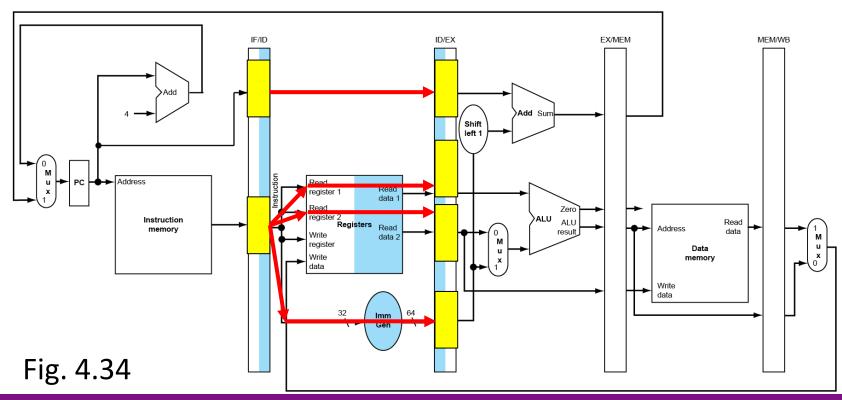
• IF/ID ← mem[PC]; IF/ID ← PC; PC ← PC + 4



ID Stage for Id

ID/EX ← IF/ID[PC]; ID/EX[A] ← Reg[IF/ID[19-15]];
 ID/EX[B] ← Reg[IF/ID[24-20]]; ID/EX ← sign-ext(IF/ID[31-20]);

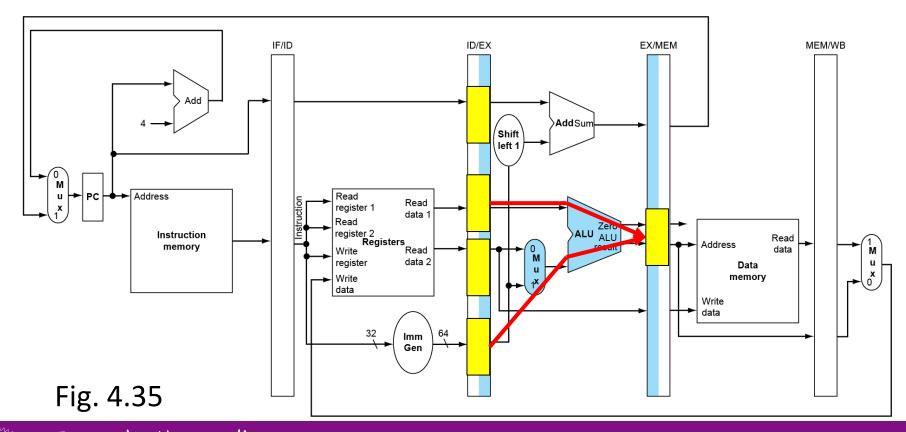
decode Instruction decode



EX Stage for Id

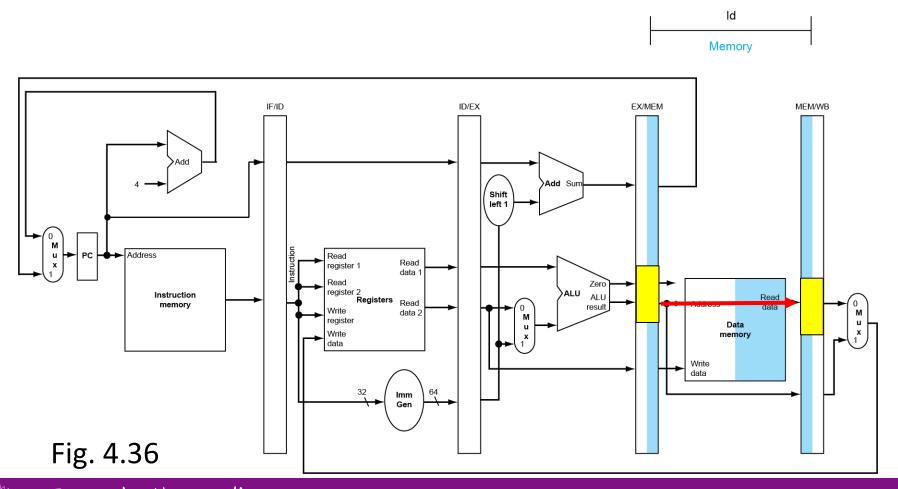
EX/MEM[ALUout] ← ID/EX[A] + ID/EX[sign-ext]





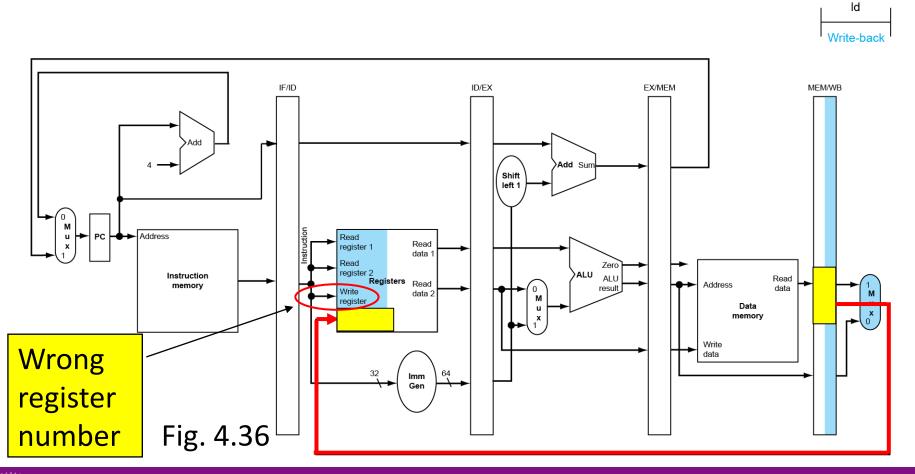
MEM Stage for Id

MEM/WB ← mem[EX/MEM[ALUout]]



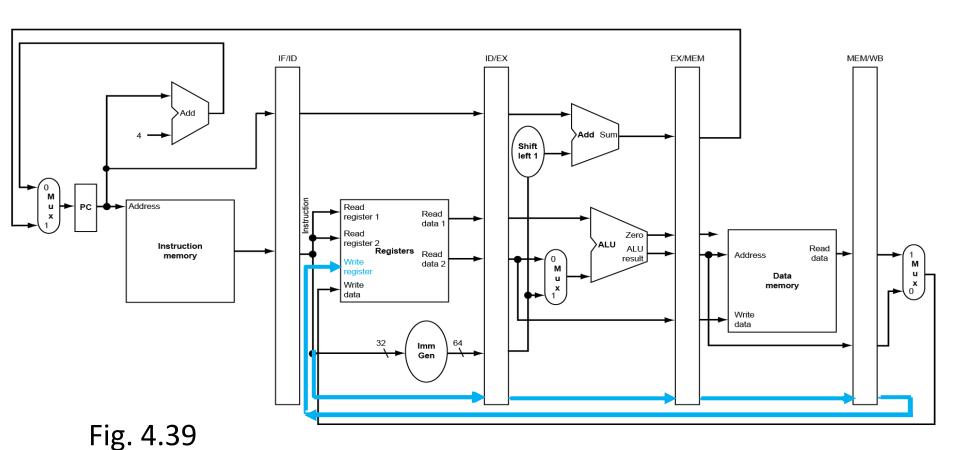
WB Stage for Id

• Reg[IF/ID[11-7]] ← MEM/WB



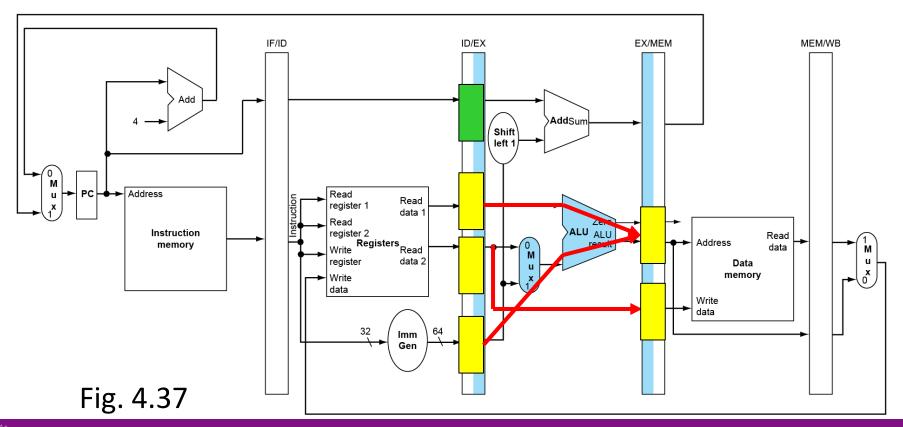
Corrected Datapath for Id

Carry rd along pipeline stages with the instruction



EX Stage for sd

EX/MEM[ALUout] ← ID/EX[A] + ID/EX[sign-ext]
 EX/MEM ← ID/EX[B]

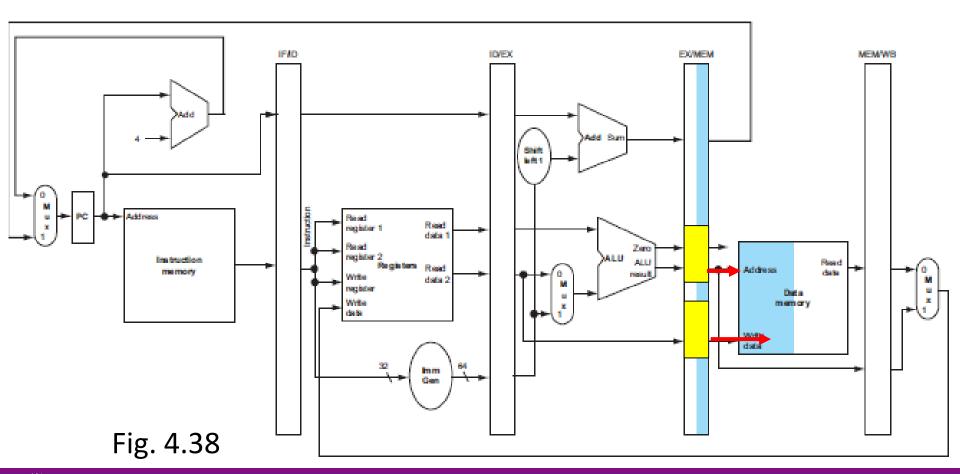


Execution

MEM Stage for sd

mem[EX/MEM[ALUout]] ← EX/MEM[B]

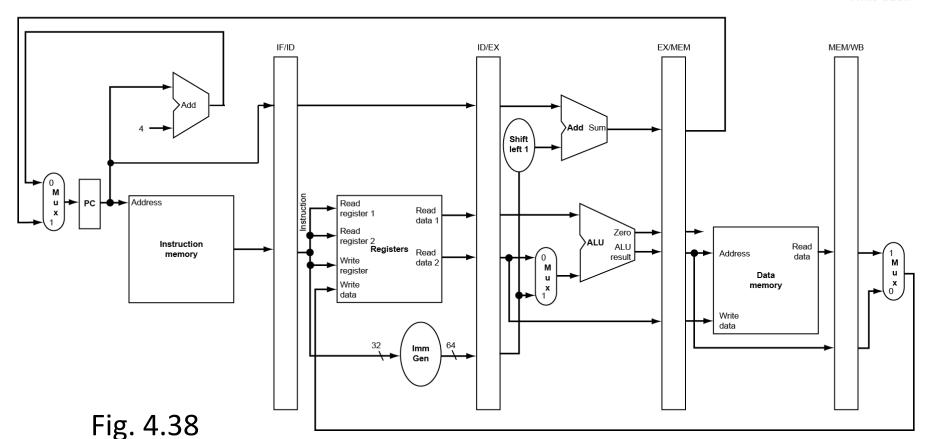
Memory



WB Stage for sd

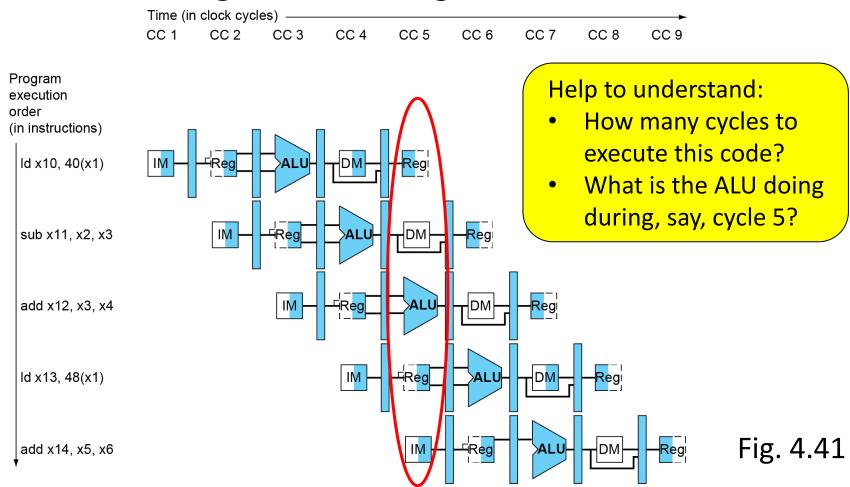
No operation

sd Write-back



Multi-Cycle Pipeline Diagram

Form showing resource usage



Multi-Cycle Pipeline Diagram

Traditional form

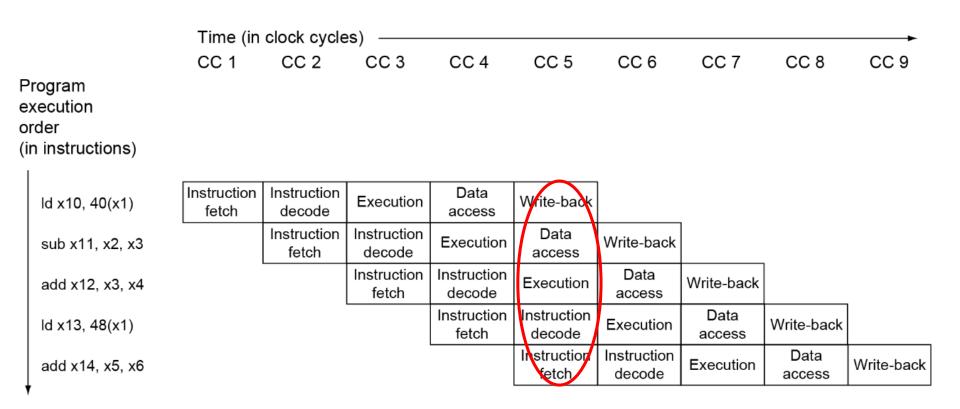
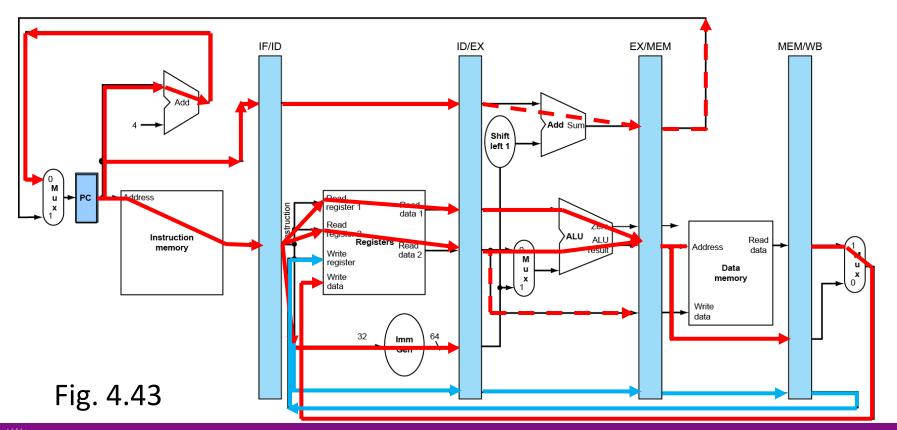


Fig. 4.42

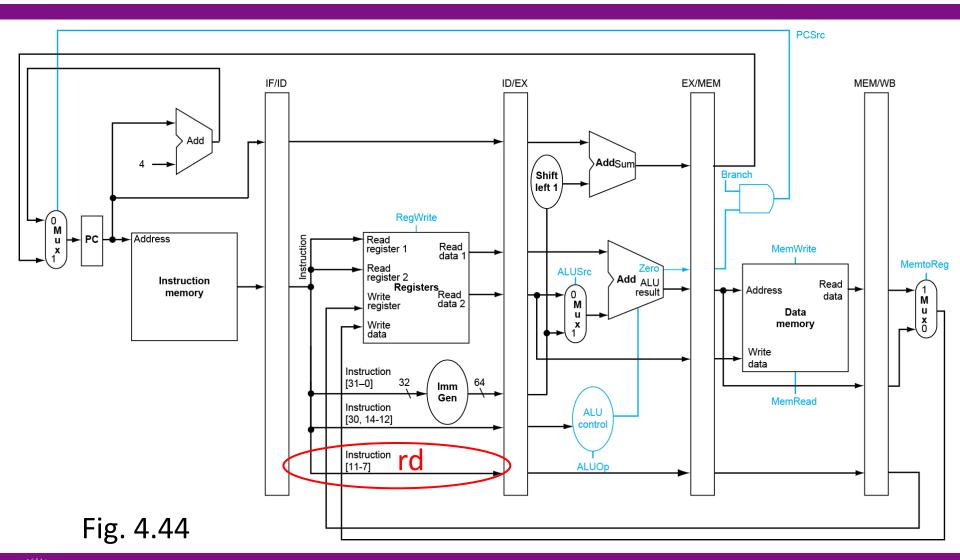
A Snapshot of the Pipeline

State of pipeline in cycle 5 of previous slide

add x14, x5, x6	ld x13, 48(x1)	add x12, x3, x4	sub x11, x2, x3	ld x10, 40(x1)	
Instruction fetch	Instruction decode	Execution	Memory	Write-back	



Pipelined Control: Control Signals



Group Signals According to Stages

Use control signals of single-cycle CPU

R-type Id sd beq

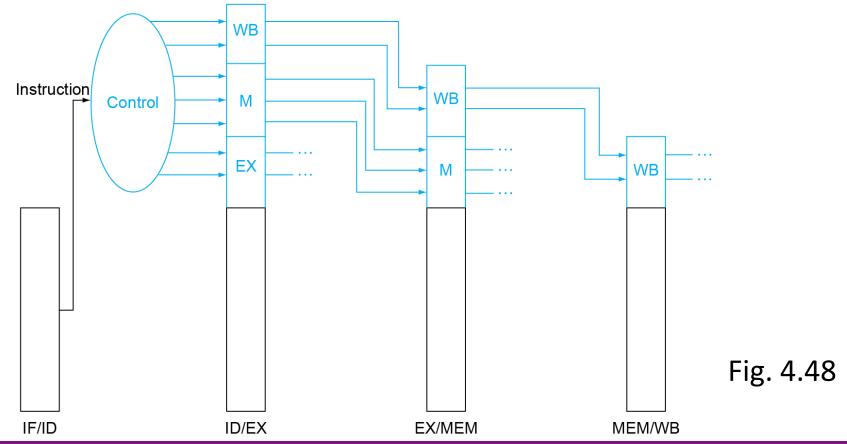
	EX Stage			ME	M Stag	WB Stage		
	ALU	ALU	ALU		Mem	Mem	Reg	Mem
	Op1	Op0	Src	Branch	Read	Write	Write	to Reg
5	1	0	0	0	0	0	1	0
	0	0	1	0	1	0	1	1
	0	0	1	0	0	1	0	X
	0	1	0	1	0	0	0	X

Fig. 4.47

 Note: There is no need to set PCsrc, because it is controlled by the "Branch" signal

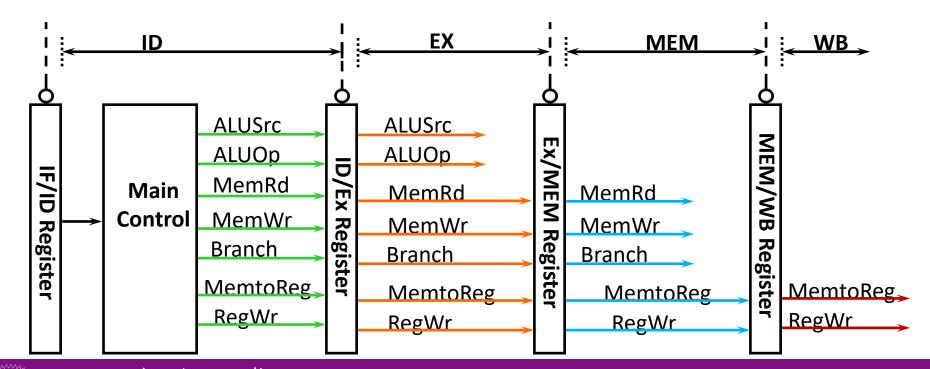
Pipelined Control

- Pass control signals along pipeline registers like data
 - Main control generates control signals during ID

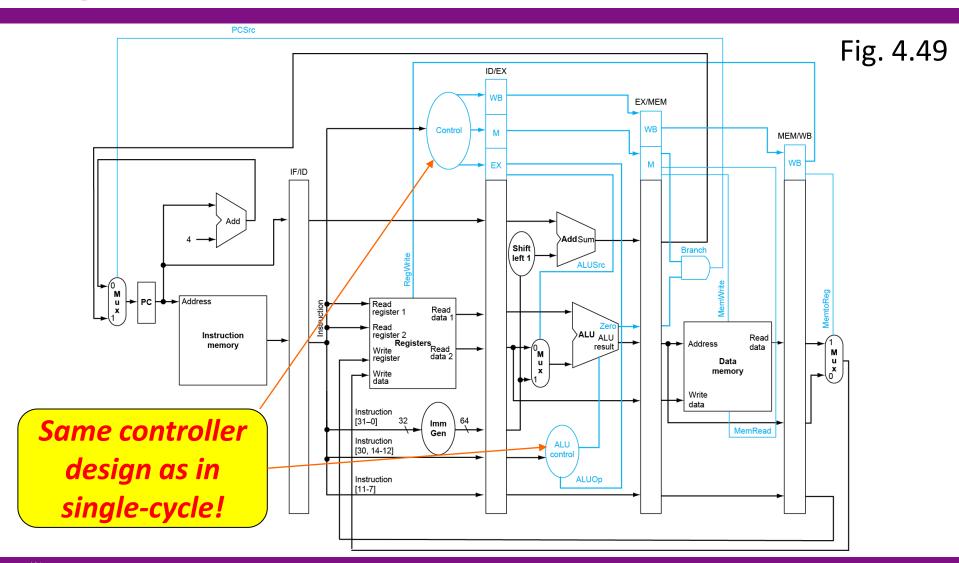


Pipelined Control

- Signals for EX (ALUSrc, ALUOp) are used 1 cycle later
- Signals for MEM (MemRd, MemWr, Branch) are used 2 cycles later
- Signals for WB (MemtoReg, RegWr) are used 3 cycles later



Datapath with Control



Summary of Pipeline Basics

- Pipelining is a fundamental concept
 - Multiple steps using distinct resources
 - Utilize capabilities of datapath by pipelined instruction processing
 - Start next instruction while working on the current one
 - Limited by length of longest stage (plus fill/flush)
- CPI = ?
- What makes it easy in RISC-V?
 - All instructions are of the same length
 - Just a few instruction formats
 - Memory operands only in loads and stores
- What makes pipelining hard? Hazards!